

# Introduction to Stan and Bayesian Inference

---

Paris Machine Learning Meetup

Dataiku User Meetup

21 September 2016

Eric Novik [enovik@stan.fit](mailto:enovik@stan.fit)



# Outline

---

- ▶ Why should you bother with Bayes
- ▶ Why should you use Stan
- ▶ Introduction to modern Bayesian workflow
- ▶ Building up a Stan model
- ▶ Brief introduction to pooling and magic of multi-level models
- ▶ Pricing books using Stan and rstanarm package
- ▶ References and guide to getting started

A high-angle, black and white photograph of a massive concrete dam. The dam's curved wall rises steeply from the bottom left towards the top right. A narrow, light-colored walkway runs along the top edge of the wall. Two small figures of people are visible on the walkway, appearing as dark silhouettes against the lighter concrete. The background is a dark, textured sky.

# Why Bayes

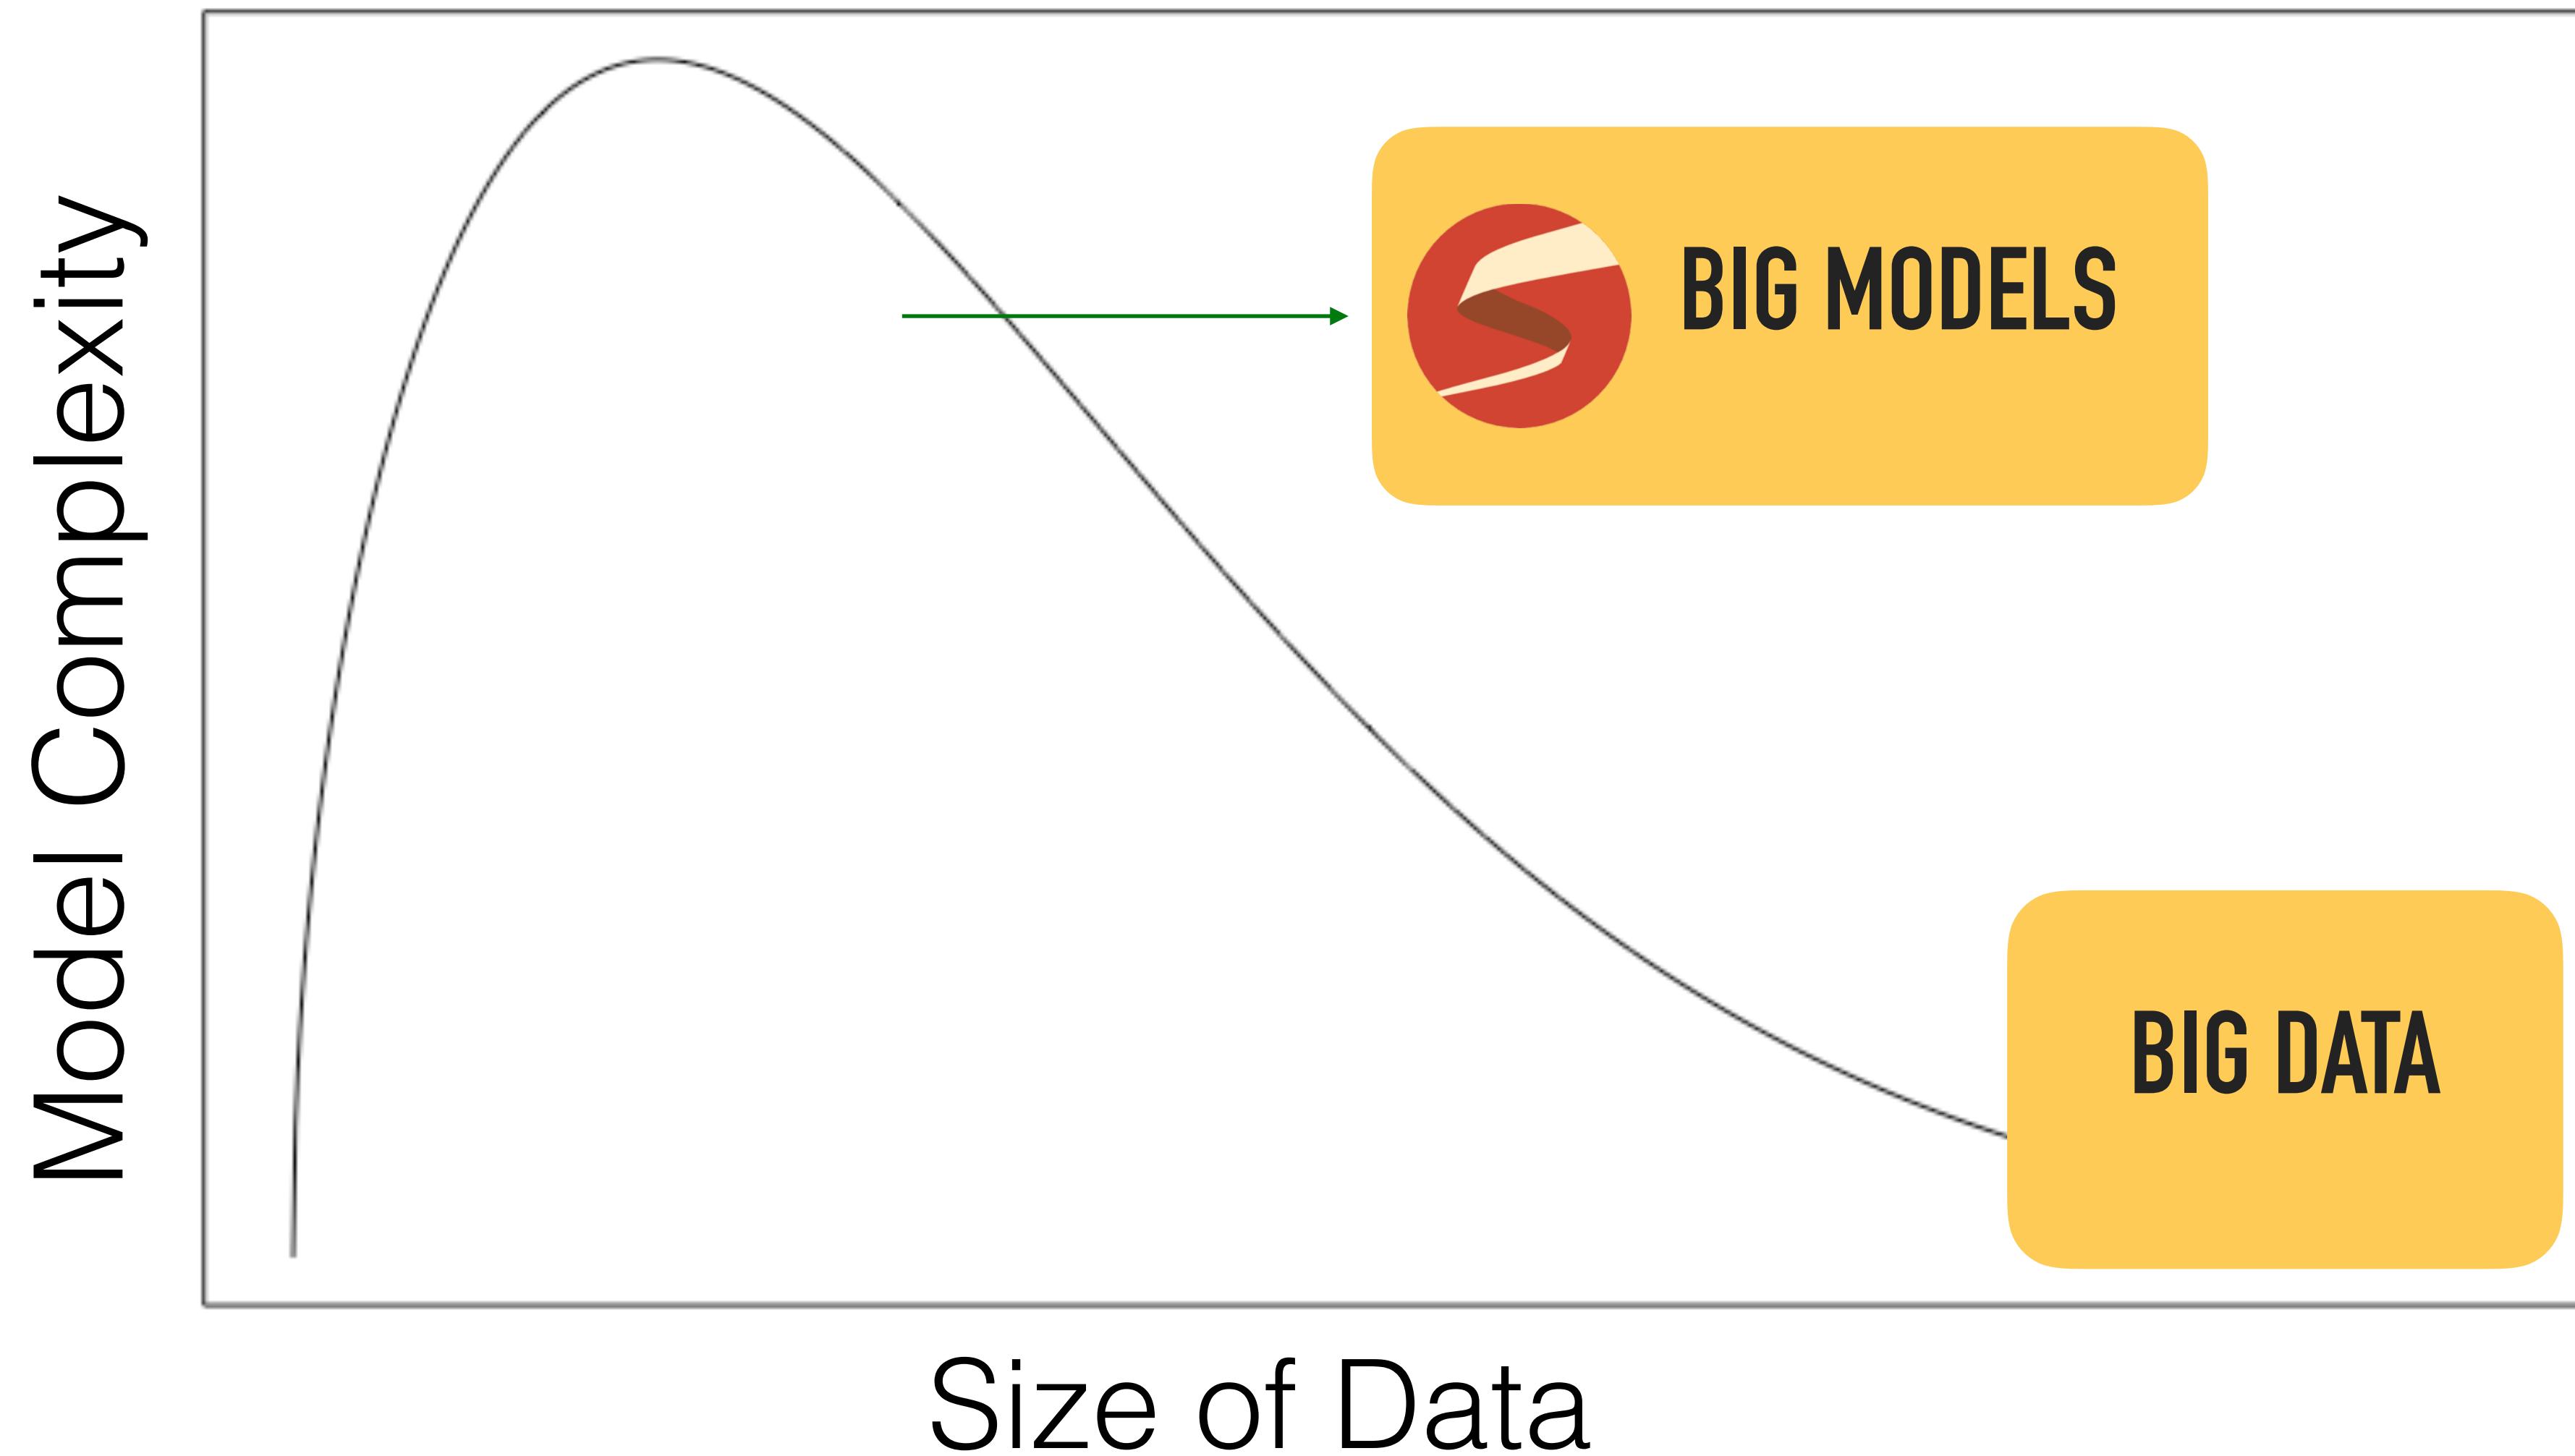
---

# Benefits of Bayesian Approach

---

- ▶ Express your beliefs about parameters **and** the data generating process
- ▶ Properly account for uncertainty at the individual and group level
- ▶ Do not collapse grouping variables (e.g. sales for of multiple products over time) and do not fit a separate model to each group
- ▶ Small data is fine if you have a strong model
- ▶ But what about Big Data?

# Big Data Need Big Models



# Traditional Machine Learning and Causal Models

---

- ▶ **Problem A:** A large retailer wants to know how many units of each product they are going to sell tomorrow
- ▶ **Problem B:** A large retailer wants to find a revenue maximizing price for all of their products
- ▶ **Data:** We observe quantity sold of each product of time, meta data about the products, and price variation
- ▶ **Question:** Which one needs a causal model?



# Why Stan

---

# What Is Stan

---

| What   | What For   |
|--|--|
| C++ Math/Stats Library   | Mathematical specification of models;<br>Automatic calculations of gradients     |
| Imperative Model Specification Language                            | Fast and simple way to specify complex<br>models                                 |
| Algorithm Toolbox  | Fit with full Bayes, approximate Bayes,<br>optimization (HMC NUTS, ADVI, L-BFGS) |
| Interfaces (Command Line, R, Python,<br>Julia, Matlab, Stata, ...) | Work in the language of your choice  |
| Interpretation Tools (shinystan)                                   | Model criticism, algorithm evaluation  |

# Who Is Using Stan

---

- ▶ 2,000+ members on the user list
- ▶ Over 10,000 manual downloads during the new release
- ▶ Stan is used for fitting climate models, clinical drug trials, genomics and cancer biology, population dynamics, psycholinguistics, social networks, finance and econometrics, professional sports, publishing, recommender systems, educational testing, and many more.



Penguin  
Random  
House



YouGov®  
What the world thinks

# Stan vs Traditional Machine Learning

- ▶ Model is directly expressed in Stan
- ▶ When in MCMC mode Stan produces draws from posterior distribution, not point estimates (MLE) of the parameters
- ▶ Fit complex models with millions of parameters
- ▶ Express and fit hierarchical models

## TRADITIONAL MACHINE LEARNING

*Model and Fitting Algorithm are Conflated and Black Box*

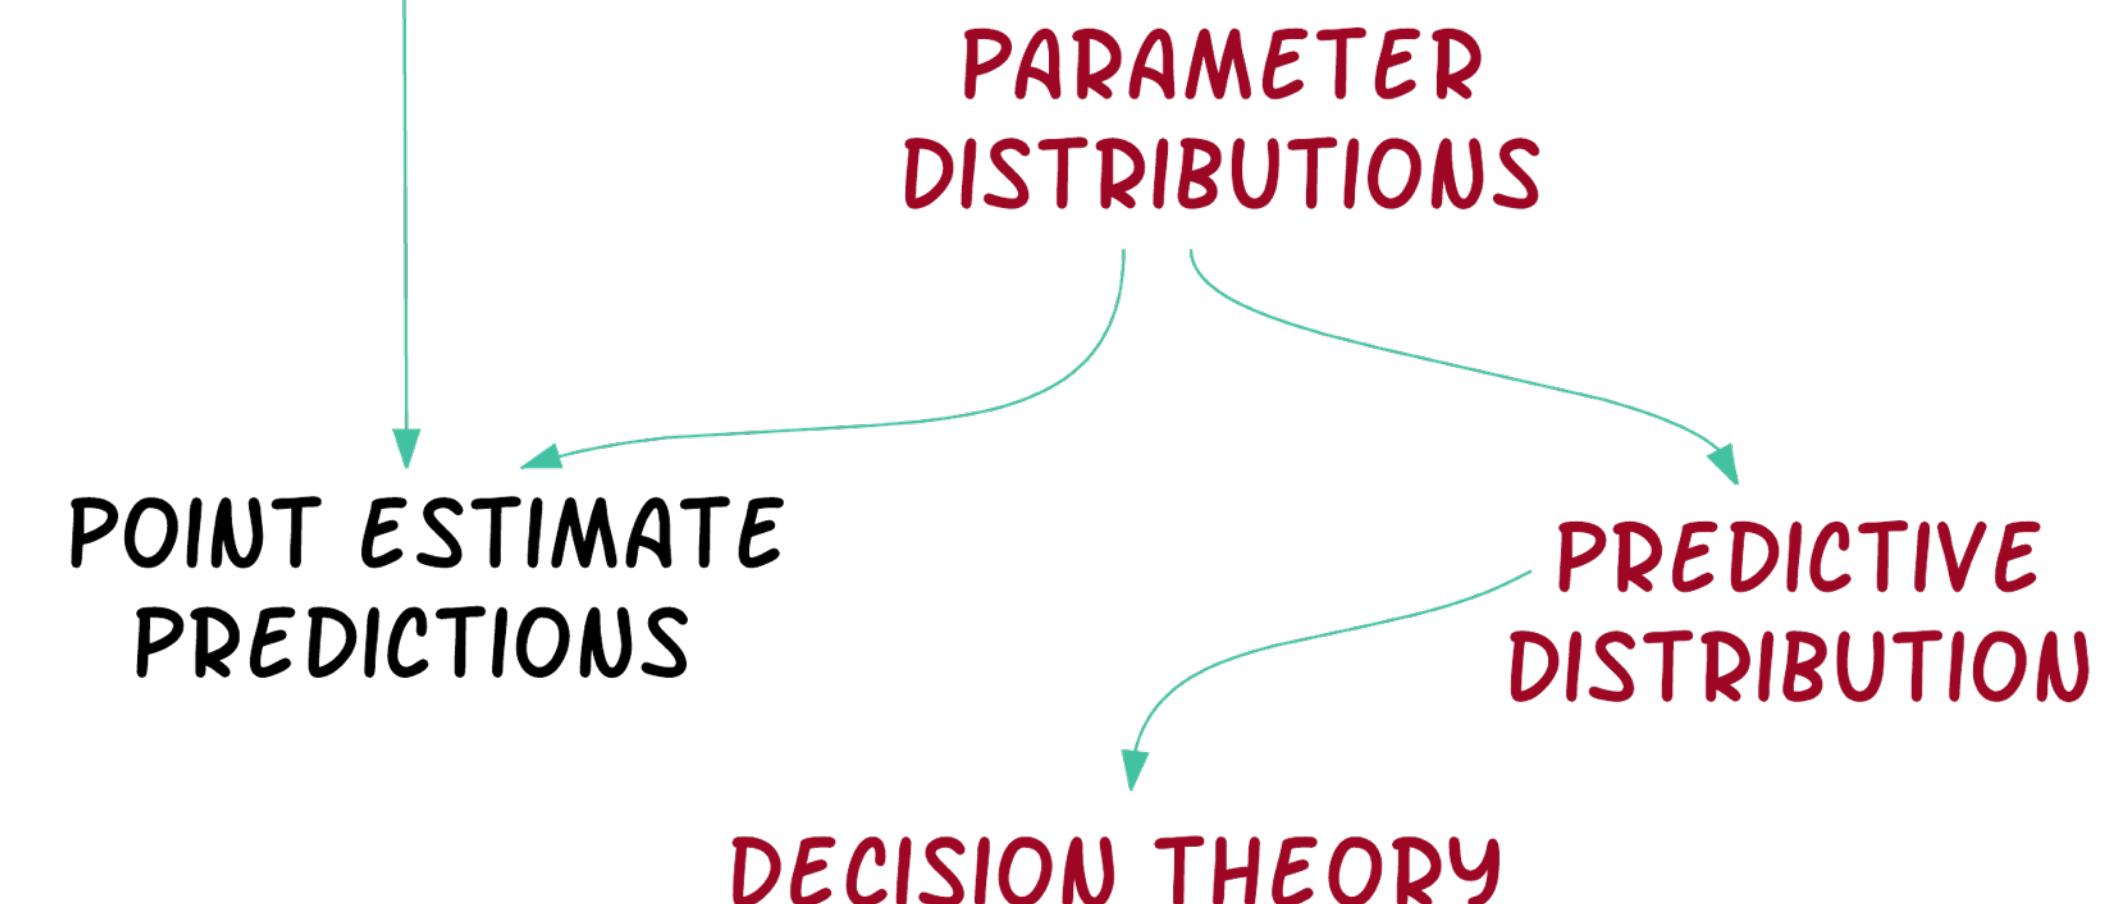
e.g. `fit = nnet(x, y, size = 2, decay = 5e-4, maxit = 200)`



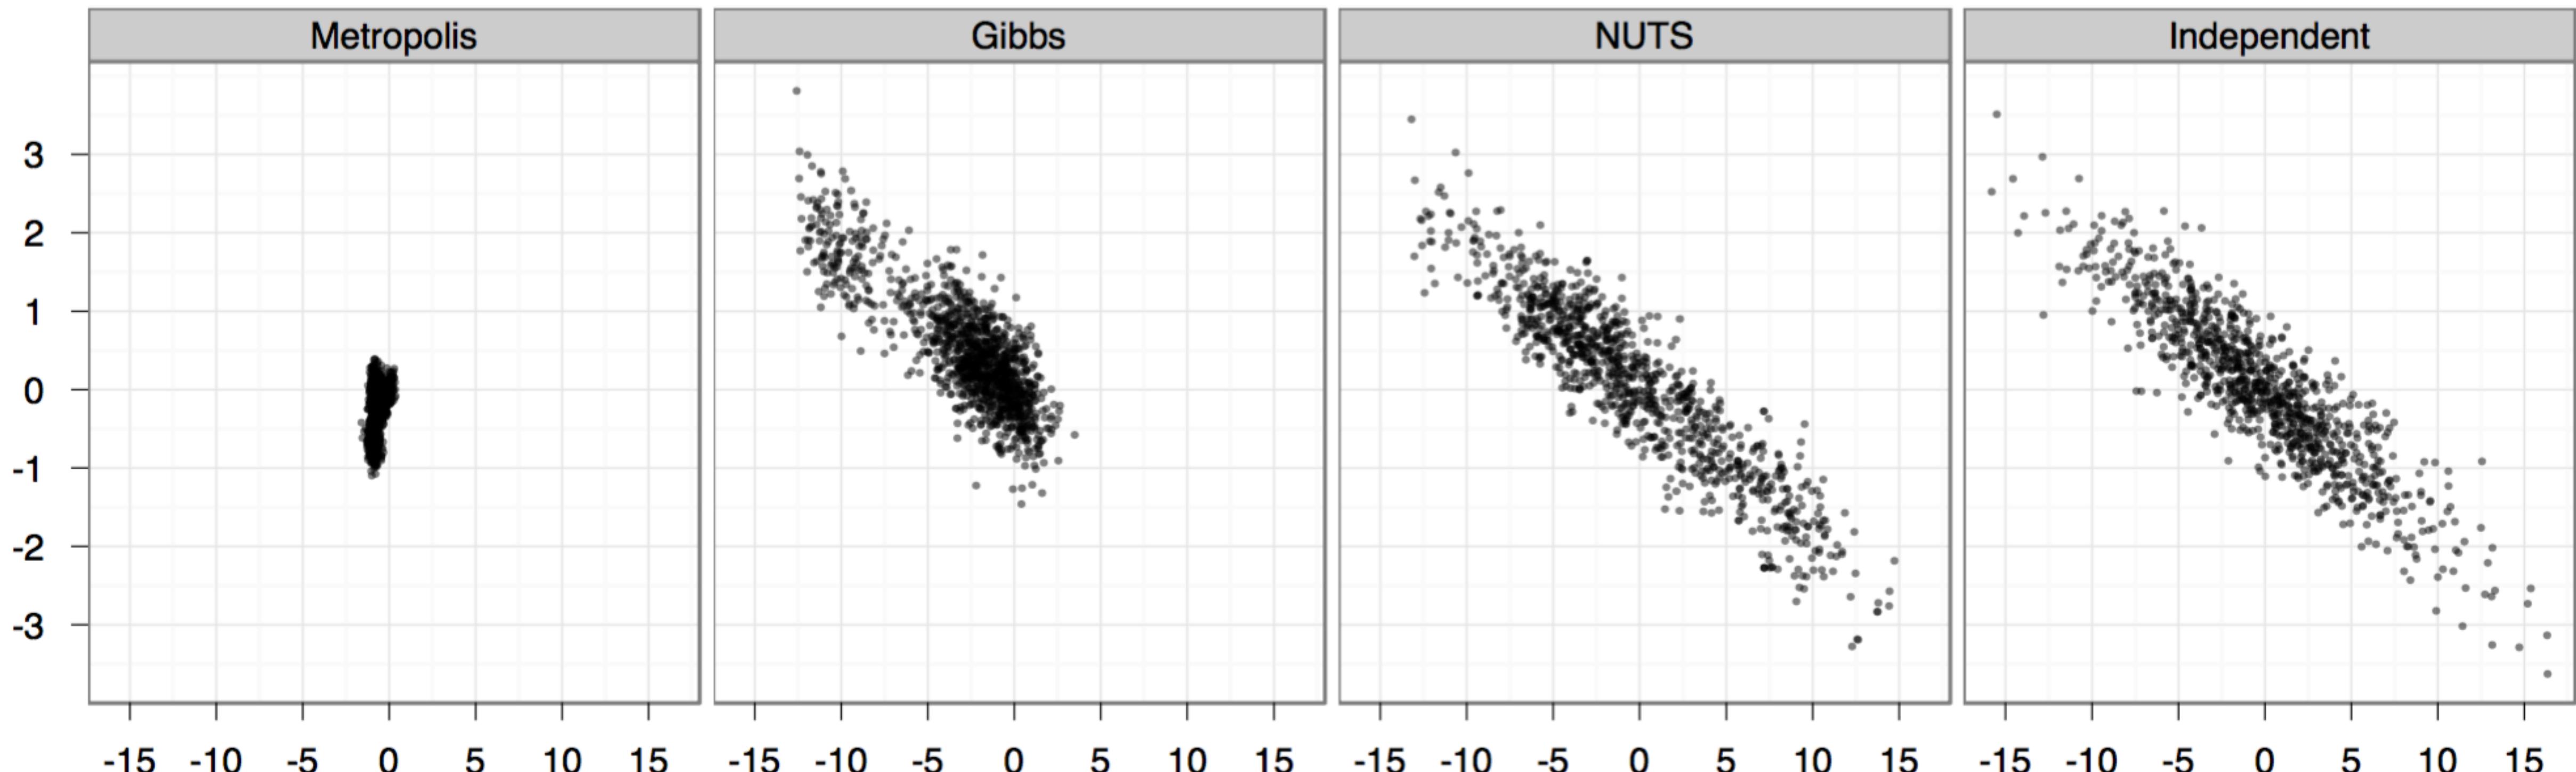
*Model is Exposed in the Stan Program*

```
model {  
    y ~ normal(alpha +  
               beta * x, sigma);  
}
```

*General Purpose Estimation Algorithms: HMC with NUTS, ADVI*



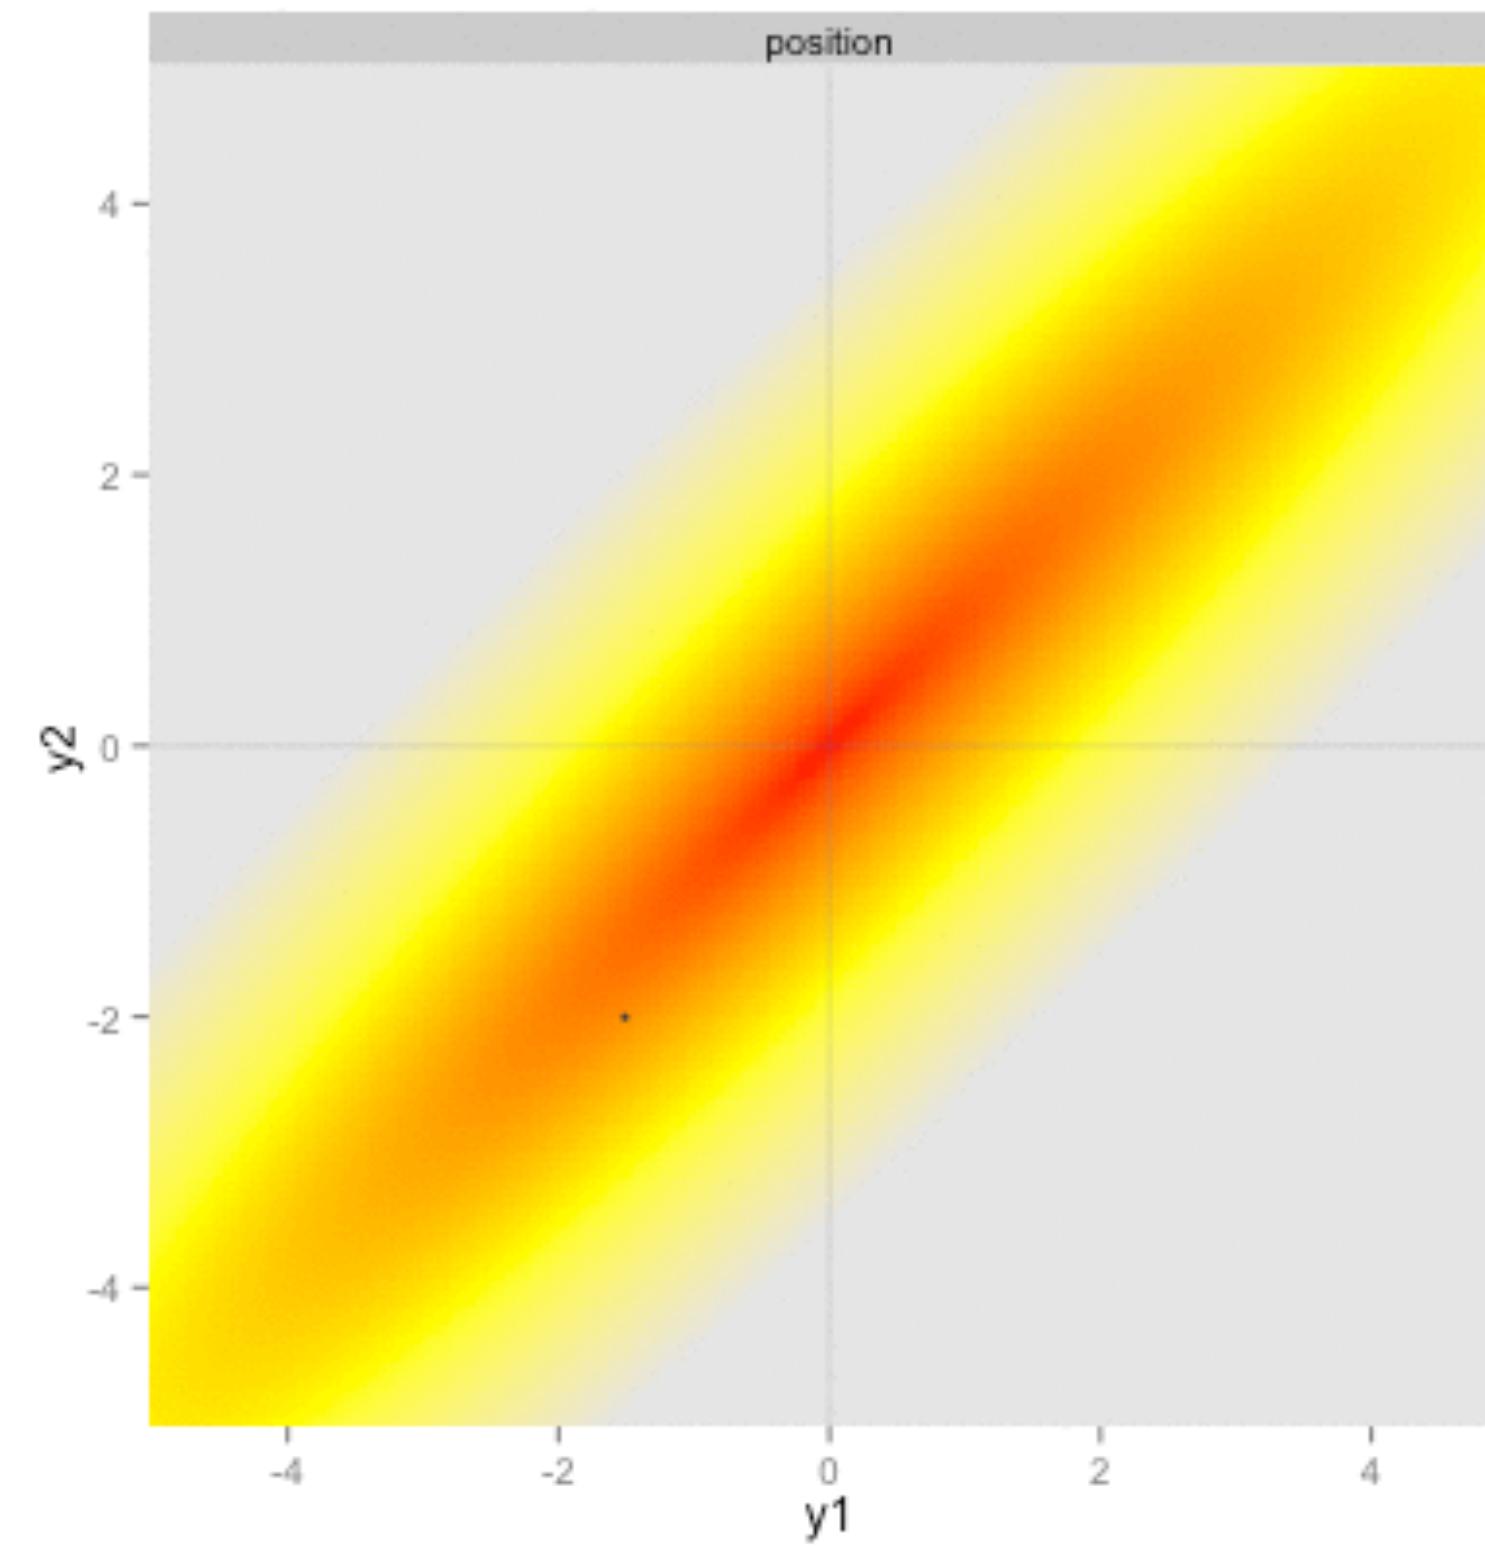
# Stan vs Gibbs and Metropolis



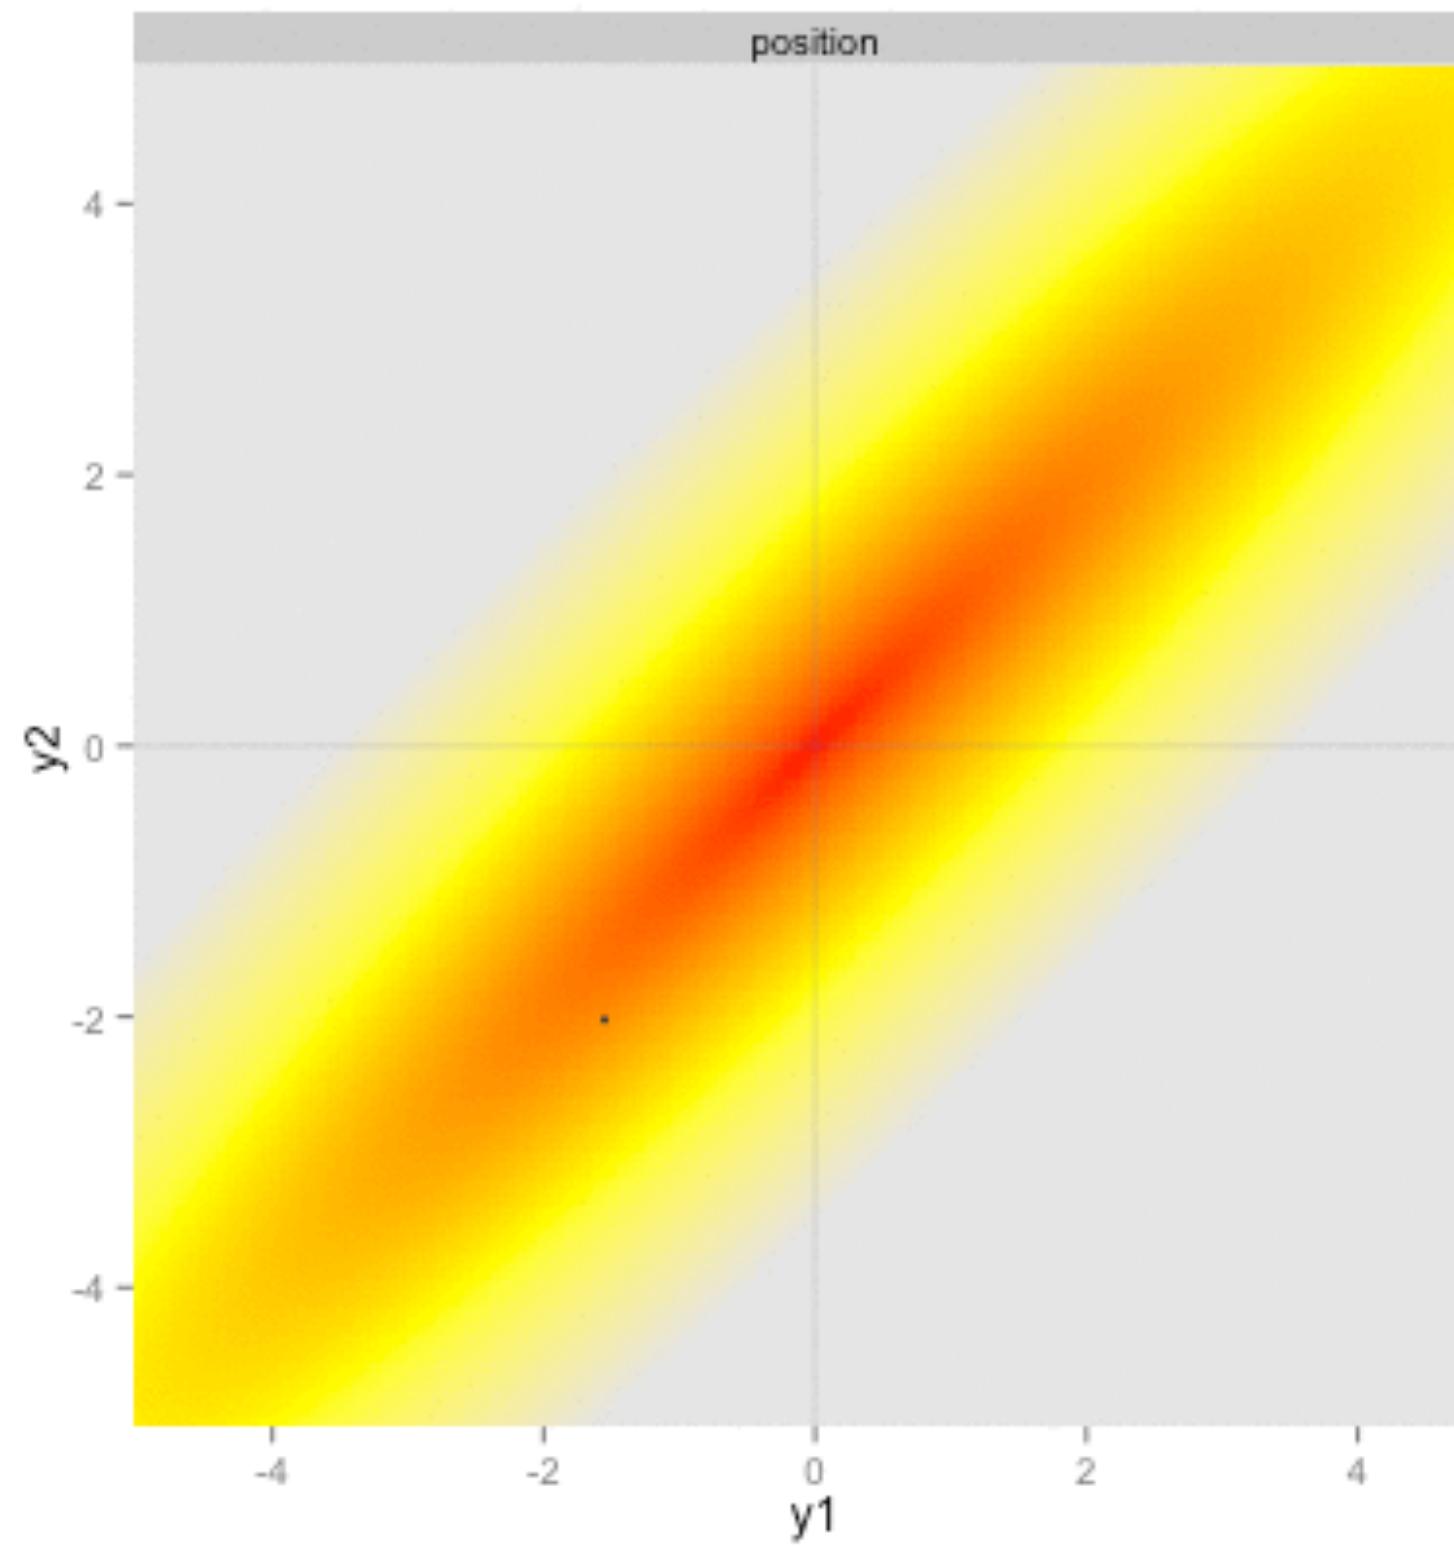
- ▶ 2-d projection of a highly correlated 250-d distribution
- ▶ 1M samples from Metropolis and 1M samples from Gibbs
- ▶ 1K samples from NUTS

# Hamiltonian Simulation

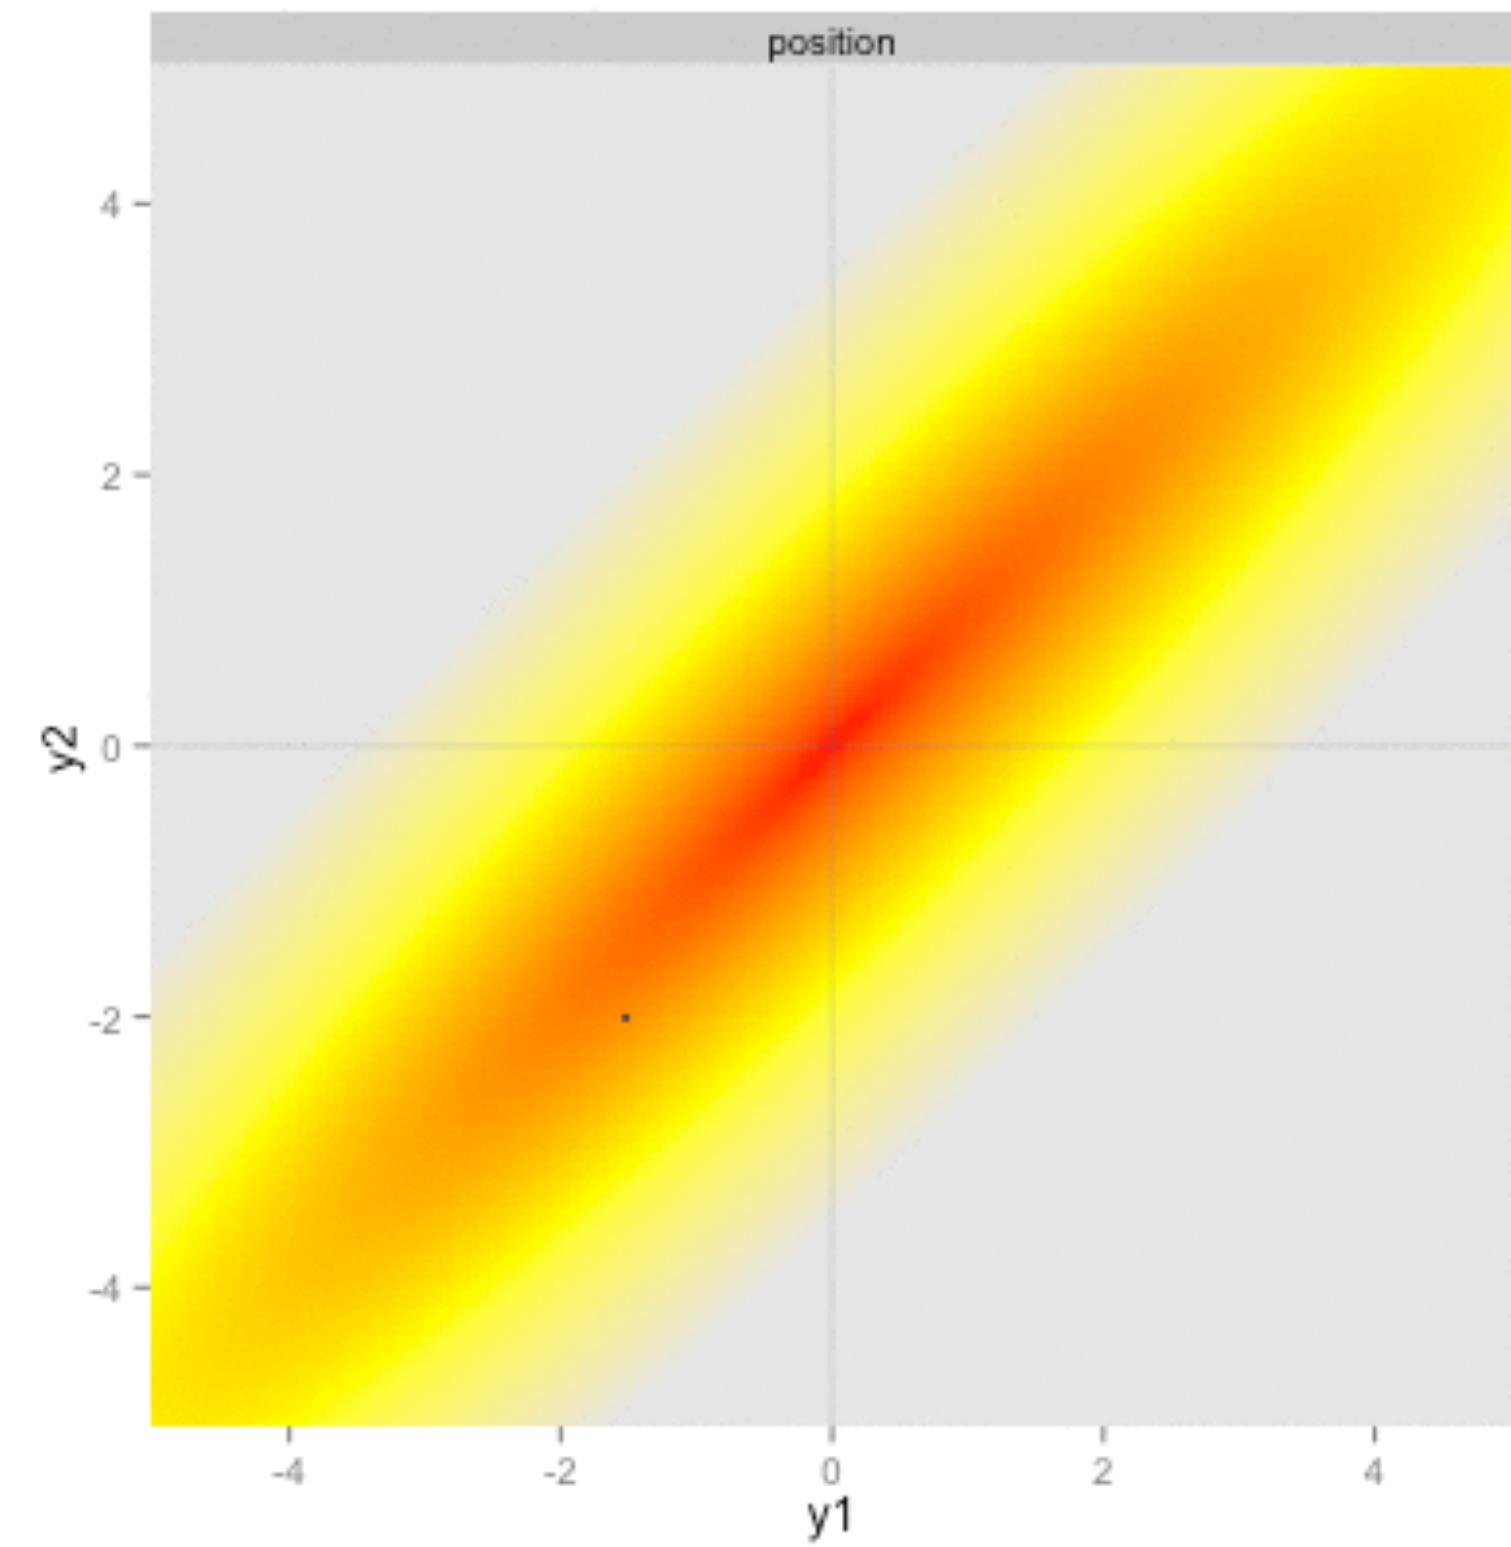
HAMILTONIAN SIMULATION  
bivariate normal ( $\rho = 0.95$ ,  $\sigma=1$ )  
init position: (-1.5, -2) init momentum: (-2, -1) stepsize: 0.005



HAMILTONIAN SIMULATION  
bivariate normal ( $\rho = 0.95$ ,  $\sigma=1$ )  
init position: (-1.5, -2) init momentum: (-2, -1) stepsize: 0.025



HAMILTONIAN SIMULATION  
bivariate normal ( $\rho = 0.95$ ,  $\sigma=1$ )  
init position: (-1.5, -2) init momentum: (-2, -1) stepsize: 0.01

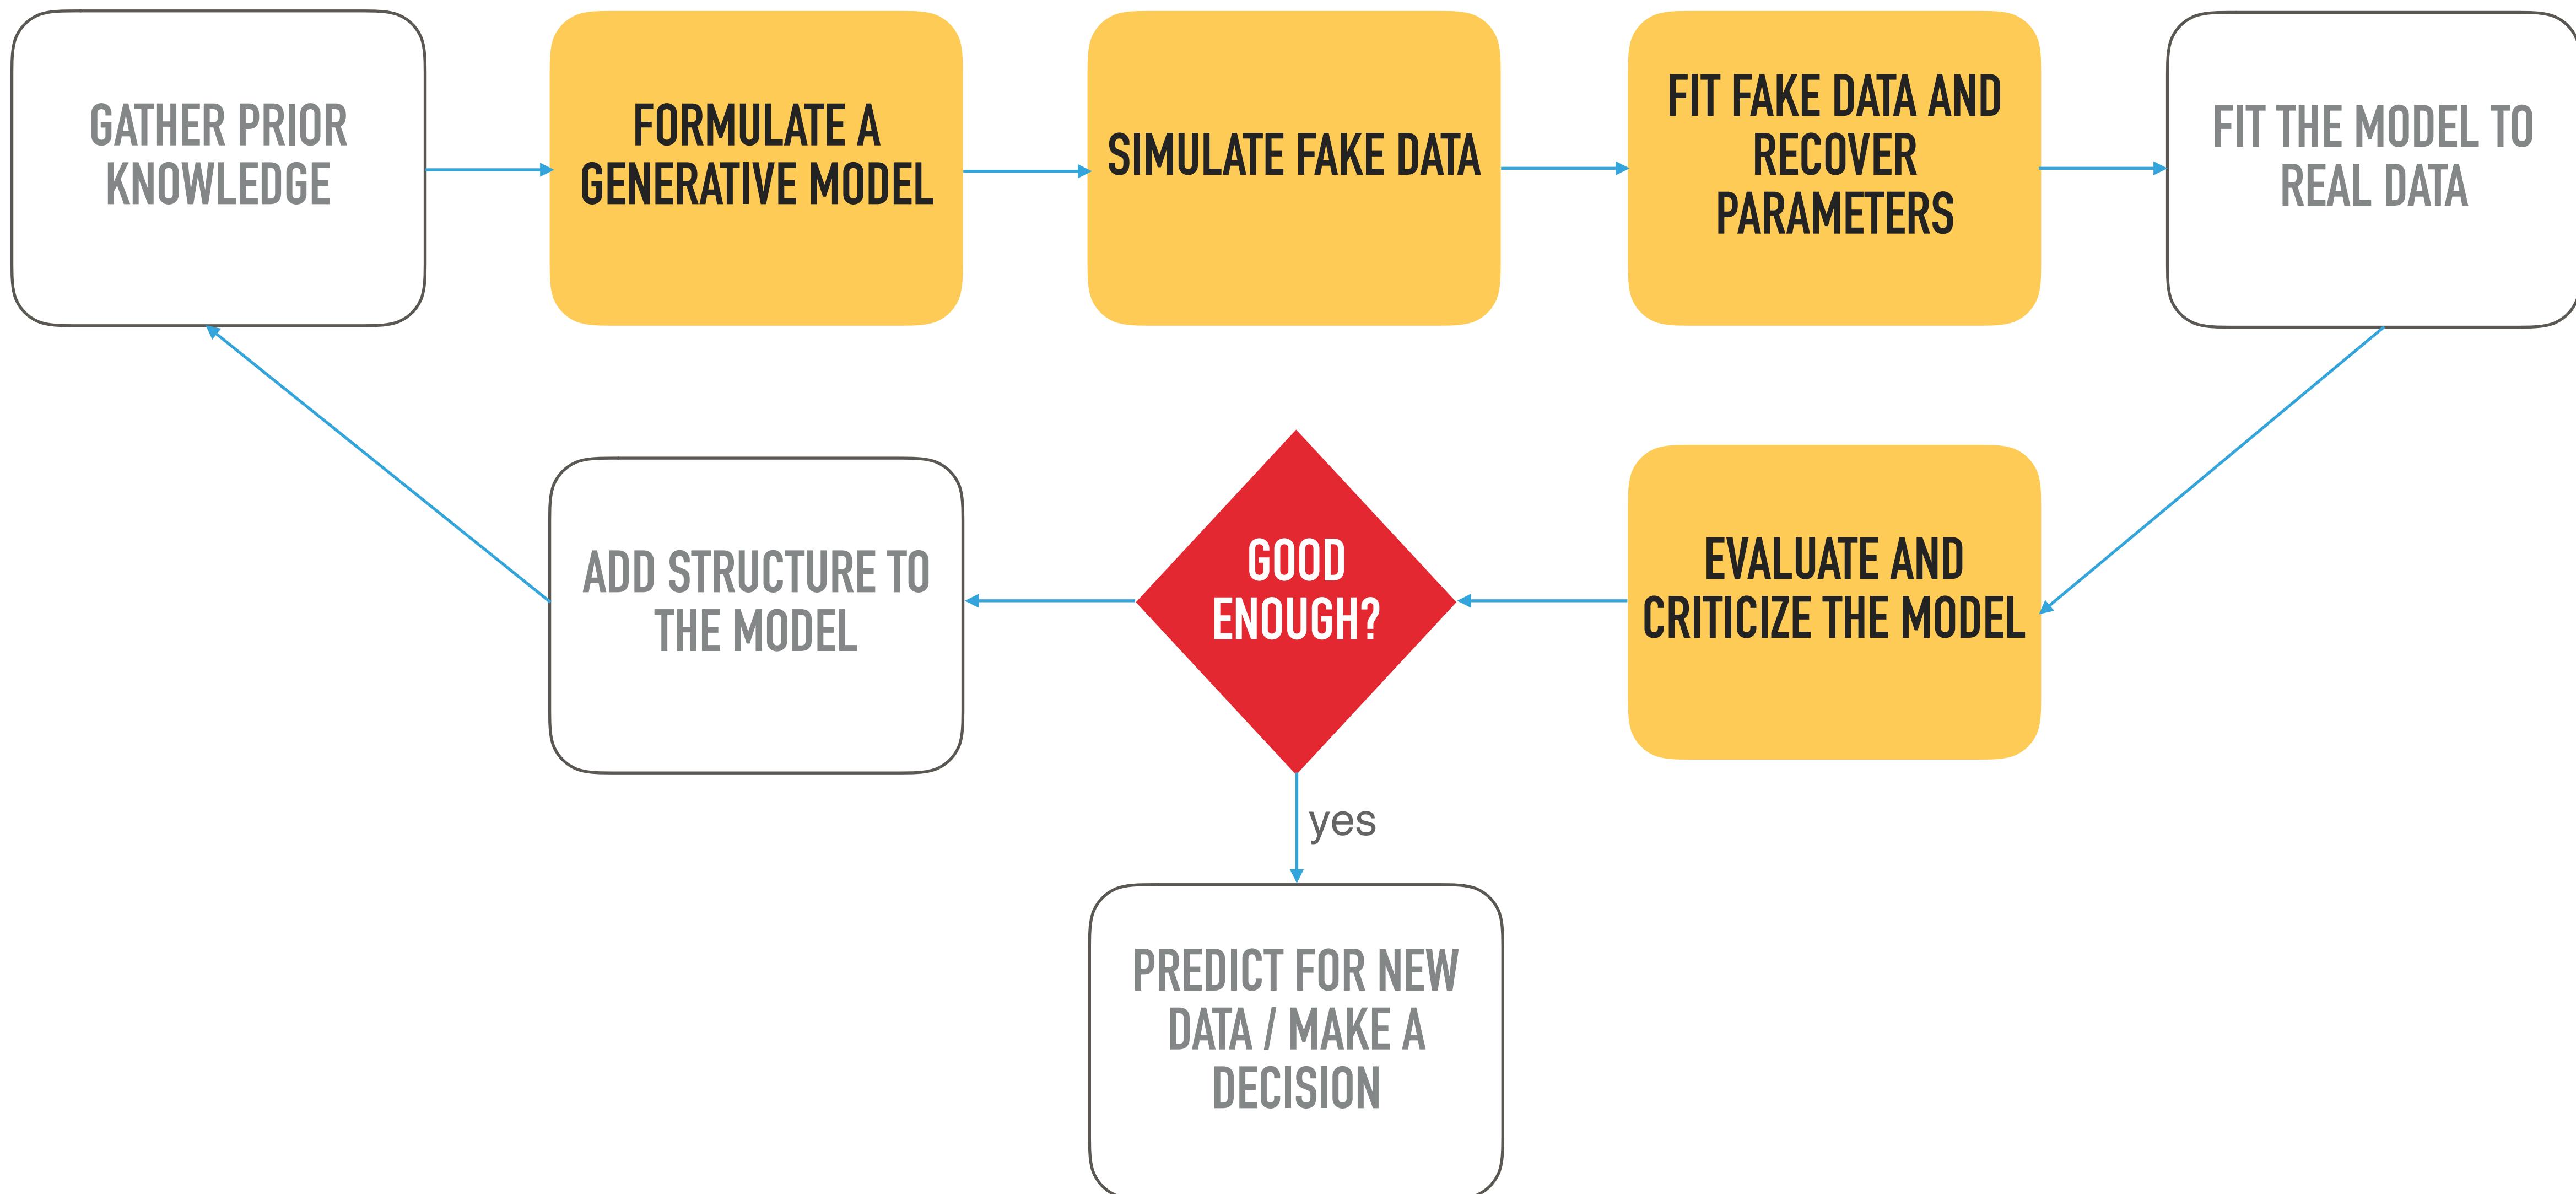


# Intro to Bayes

---

with Modern Bayesian Workflow

# Bayesian Workflow



# Bayesian Machinery

---

- The joint probability of data **y** and unknown parameter **theta**:

$$p(y, \theta) = p(y|\theta) * p(\theta)$$

$$p(y, \theta) = p(\theta|y) * p(y)$$

- The conditional probability of **theta** given **y**:

$$p(\theta|y) = \frac{p(y|\theta) * p(\theta)}{p(y)} = \frac{p(y|\theta) * p(\theta)}{\int p(y, \theta) d\theta} = \frac{p(y|\theta) * p(\theta)}{\int p(y|\theta) * p(\theta) d\theta}$$

Likelihood    Prior  
Marginal Likelihood

$$\propto p(y|\theta) * p(\theta)$$

# Bernoulli Model

- If we model each occurrence as independent, the joint model can be written as:

$$p(y, \theta) = \prod_{n=1}^N \theta^{y_n} * (1 - \theta)^{1-y_n} = \theta^{\sum_{n=1}^N y_n} * (1 - \theta)^{\sum_{n=1}^N (1-y_n)}$$

Bernoulli Likelihood  $p(y|\theta)$

- What happened to the prior,  $p(\theta)$
- On the log scale:

$$\log(p(y, \theta)) = \sum_{n=1}^N y_n * \log(\theta) + \sum_{n=1}^N (1 - y_n) * \log(1 - \theta)$$

```
data <- list(N = 5,
              y = c(0, 1, 1, 0, 1))

# log probability function
lp <- function(theta, data) {
  lp <- 0
  for (i in 1:data$N) {
    lp <- lp + log(theta) * data$y[i] +
      log(1 - theta) * (1 - data$y[i])
  }
  return(lp)
}
```

# Bernoulli Model

---

```
# using dbinom()
lp_dbinom <- function(theta, d) {
  lp <- 0
  for (i in 1:length(theta))
    lp[i] <- sum(dbinom(d$y, size = 1,
                         prob = theta[i],
                         log = TRUE))
  return(lp)
}

> lp(c(0.6, 0.7), data)
[1] -3.365058 -3.477970

> lp_dbinom(c(0.6, 0.7), data)
[1] -3.365058 -3.477970
```

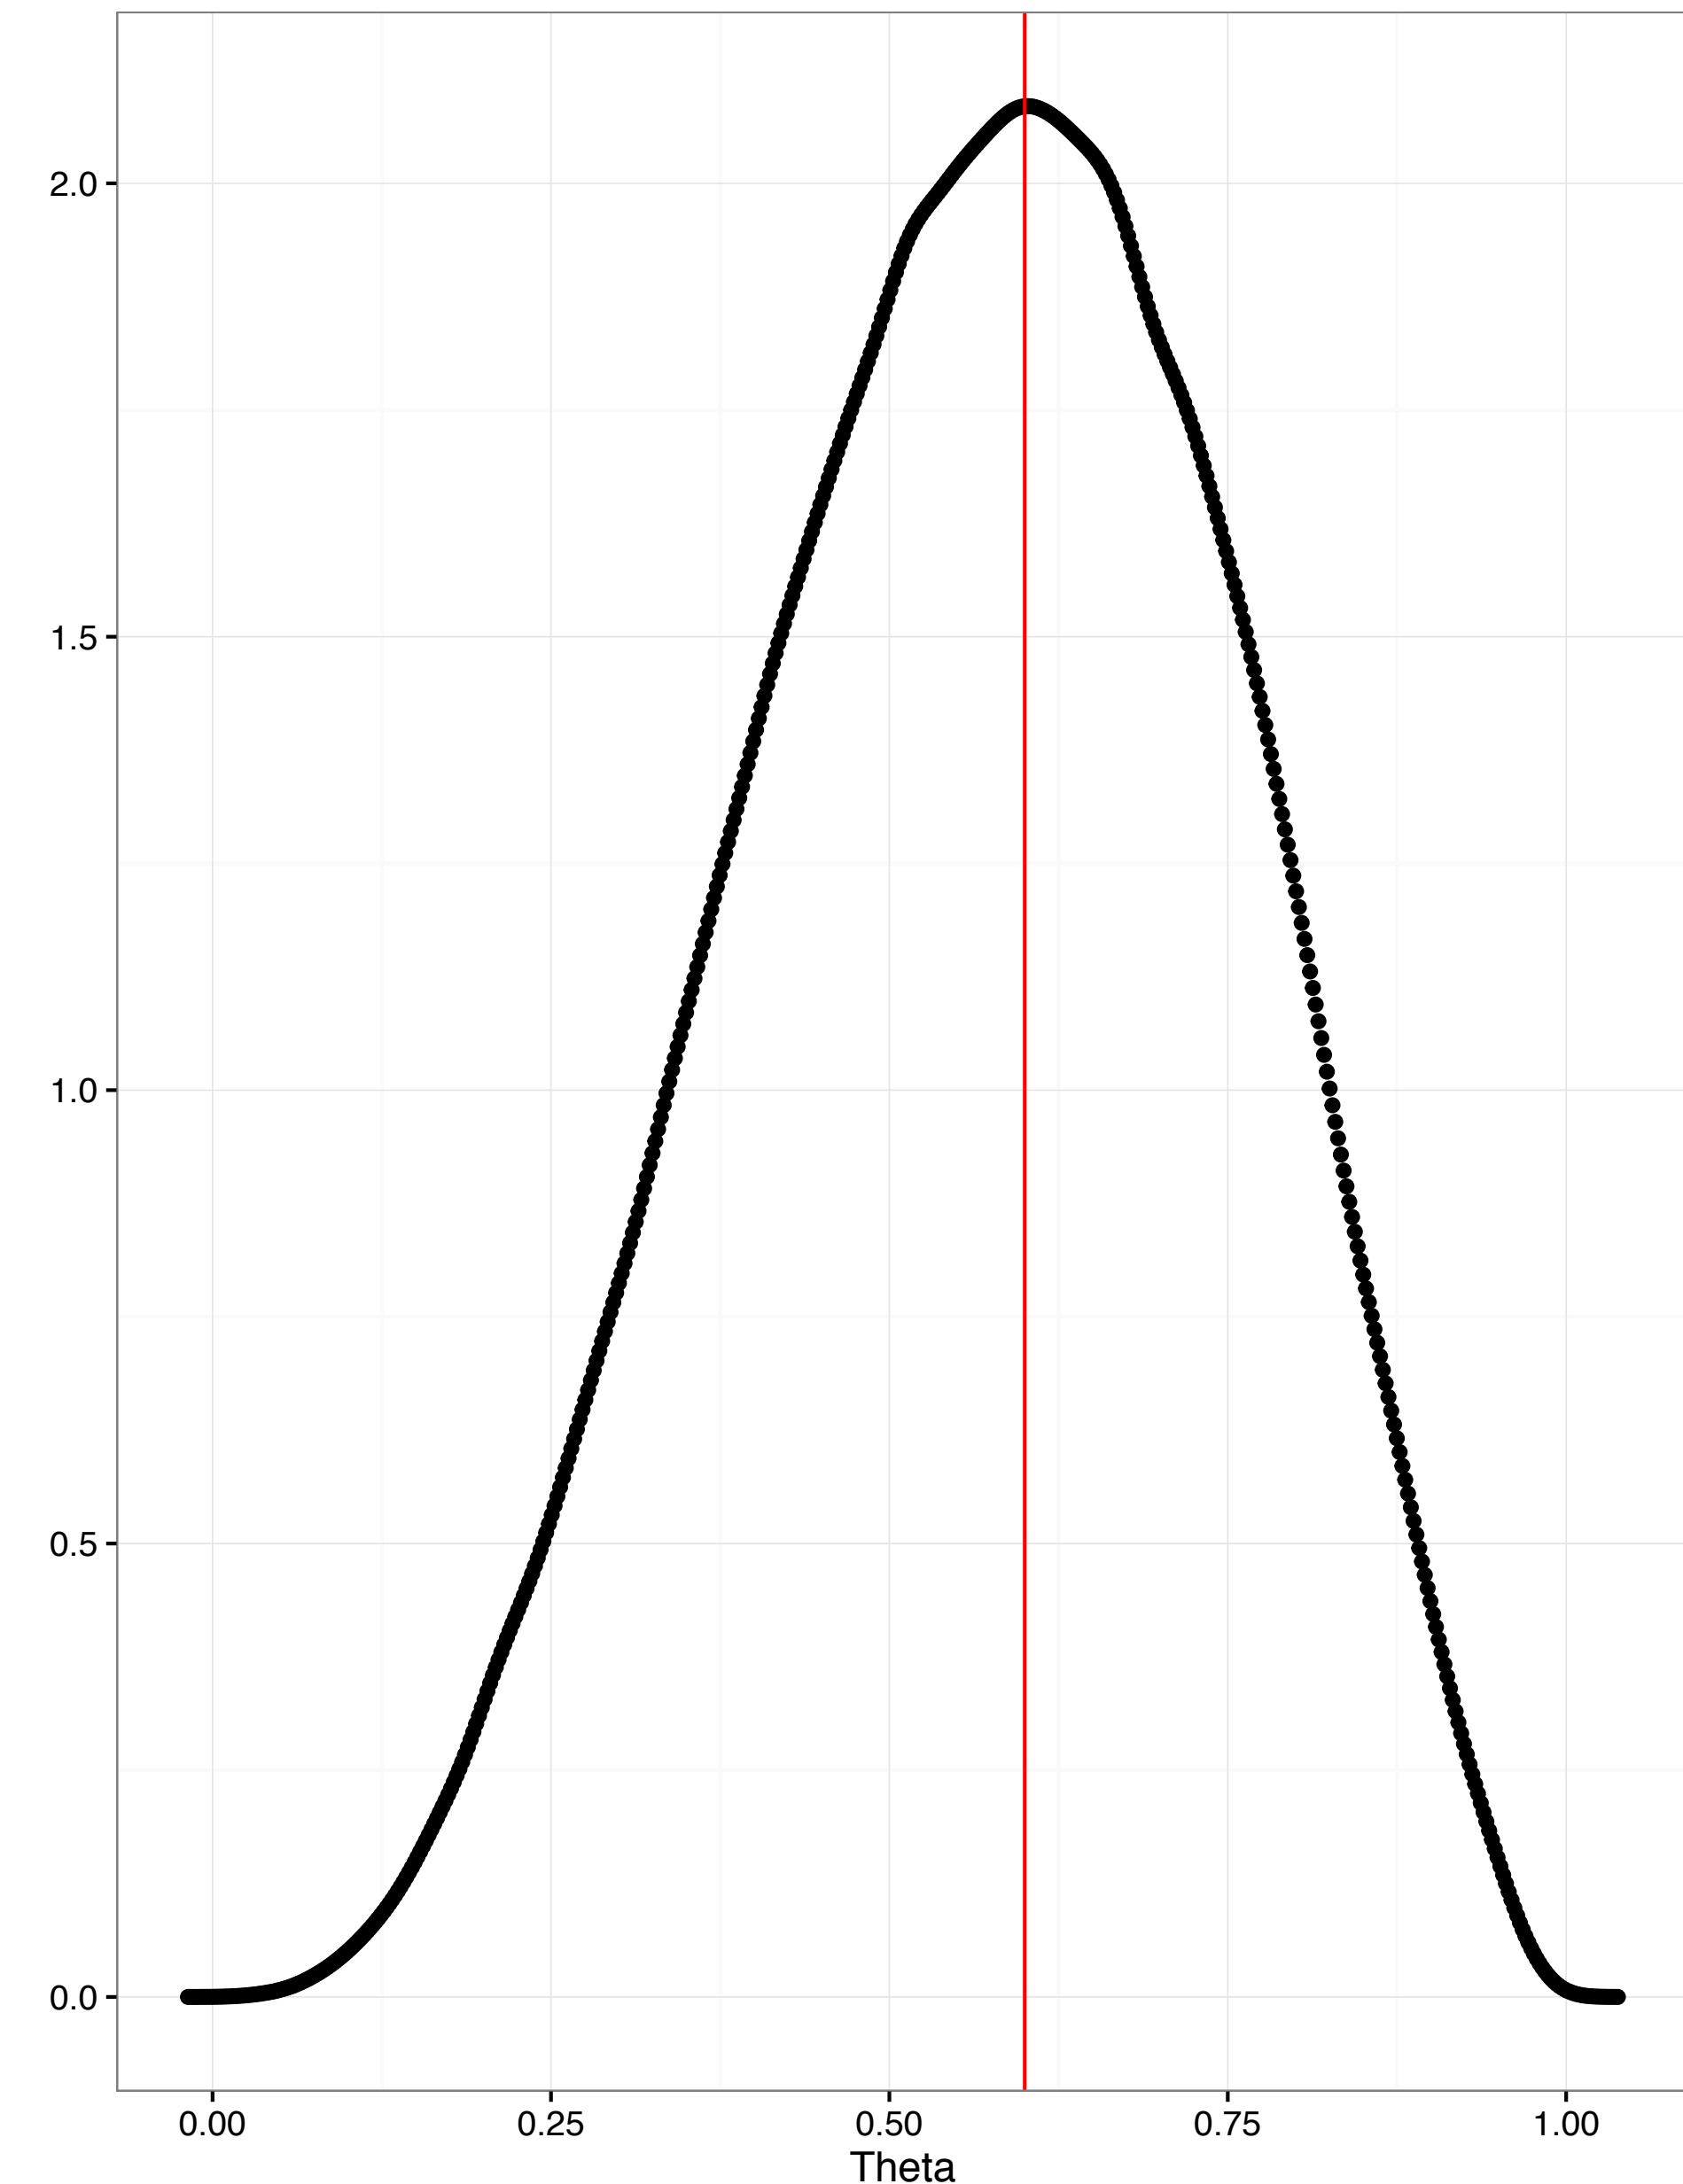
# Bernoulli Model

```
# generate the parameter grid
theta <- seq(0.001, 0.999,
             length.out = 250)

# log p(theta | y)
posterior <- lp(theta = theta, data)
posterior <- exp(log_prob)

# normalize
posterior <- posterior / sum(posterior)

# sample from the posterior
post_draws <- sample(theta, size = 1e5,
                      replace = TRUE,
                      prob = posterior)
post_dens <- density(post_draws)
mle <- sum(data$y) / data$N
> mle
[1] 0.6
```



# Same Model in Stan

```
data {  
    int<lower=1> N;  
    int<lower=0, upper=1> y[N];  
}  
  
parameters {  
    real theta;  
}  
  
model {  
    for (n in 1:N)  
        target += y[n] * log(theta) +  
                    (1 - y[n]) * log(1 - theta);  
}
```

```
data {  
    int<lower=1> N;  
    int<lower=0, upper=1> y[N];  
}  
  
parameters {  
    real<lower=0, upper=1> theta;  
}  
  
model {  
    y ~ bernoulli(theta);  
}
```

$$\log(p(y, \theta)) = \sum_{n=1}^N y_n * \log(\theta) + \sum_{n=1}^N (1 - y_n) * \log(1 - \theta)$$

# Product Pricing

---

Basic Model and Data Simulation

# Anlytical Problem

---

- ▶ A large publisher has hundreds of thousands of books in the catalog
- ▶ Every year, thousands of new books (products) are published
- ▶ There are also new authors, repeat authors, genres, seasonality, and so on
- ▶ Publisher wants to maximize revenue but if uncertainty is high, maximize quantity sold
- ▶ How should we model this? (and what is this)?



# Basic Model for Quantity Sold

$$qty = f(price, price^2, product\_age, \dots)$$

- ▶ For a Gaussian model, and one product:

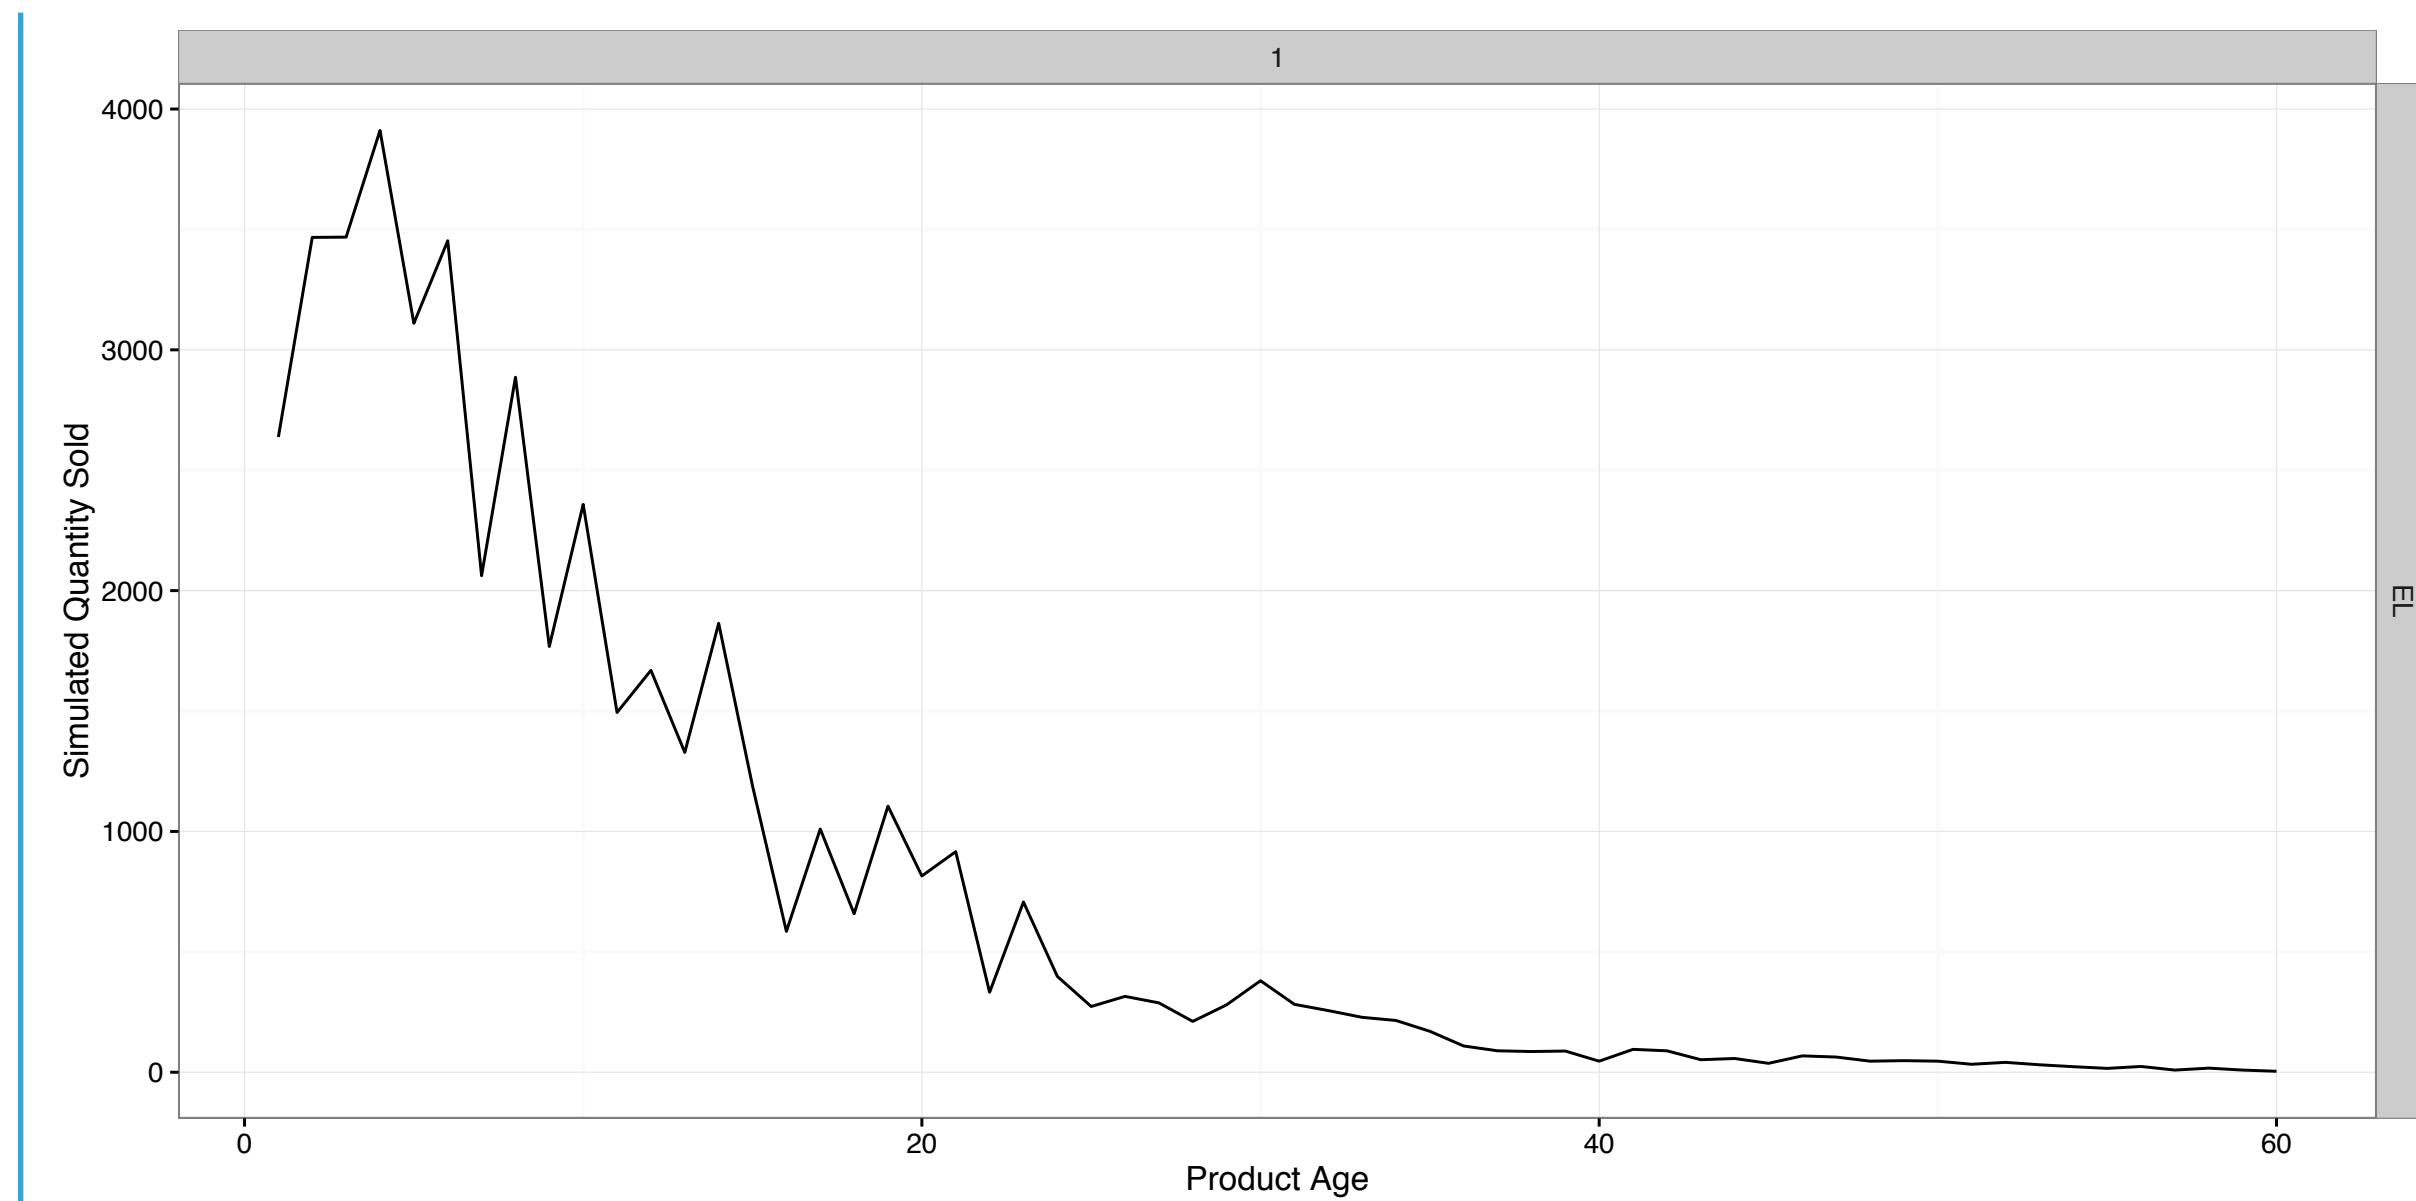
$$qty_i \sim N(X_i\beta, \sigma^2)$$

- ▶ For products that sell thousands of units we would fit a log-log model
- ▶ For lower volume products that sometimes sell zero units, we fit a count model that does not force the mean to be equal to the variance

$$qty \sim NegativeBinomial(\mu, \phi)$$

$$\mu = exp(\alpha + \beta_1 * product\_age + \beta_2 * price + \dots)$$

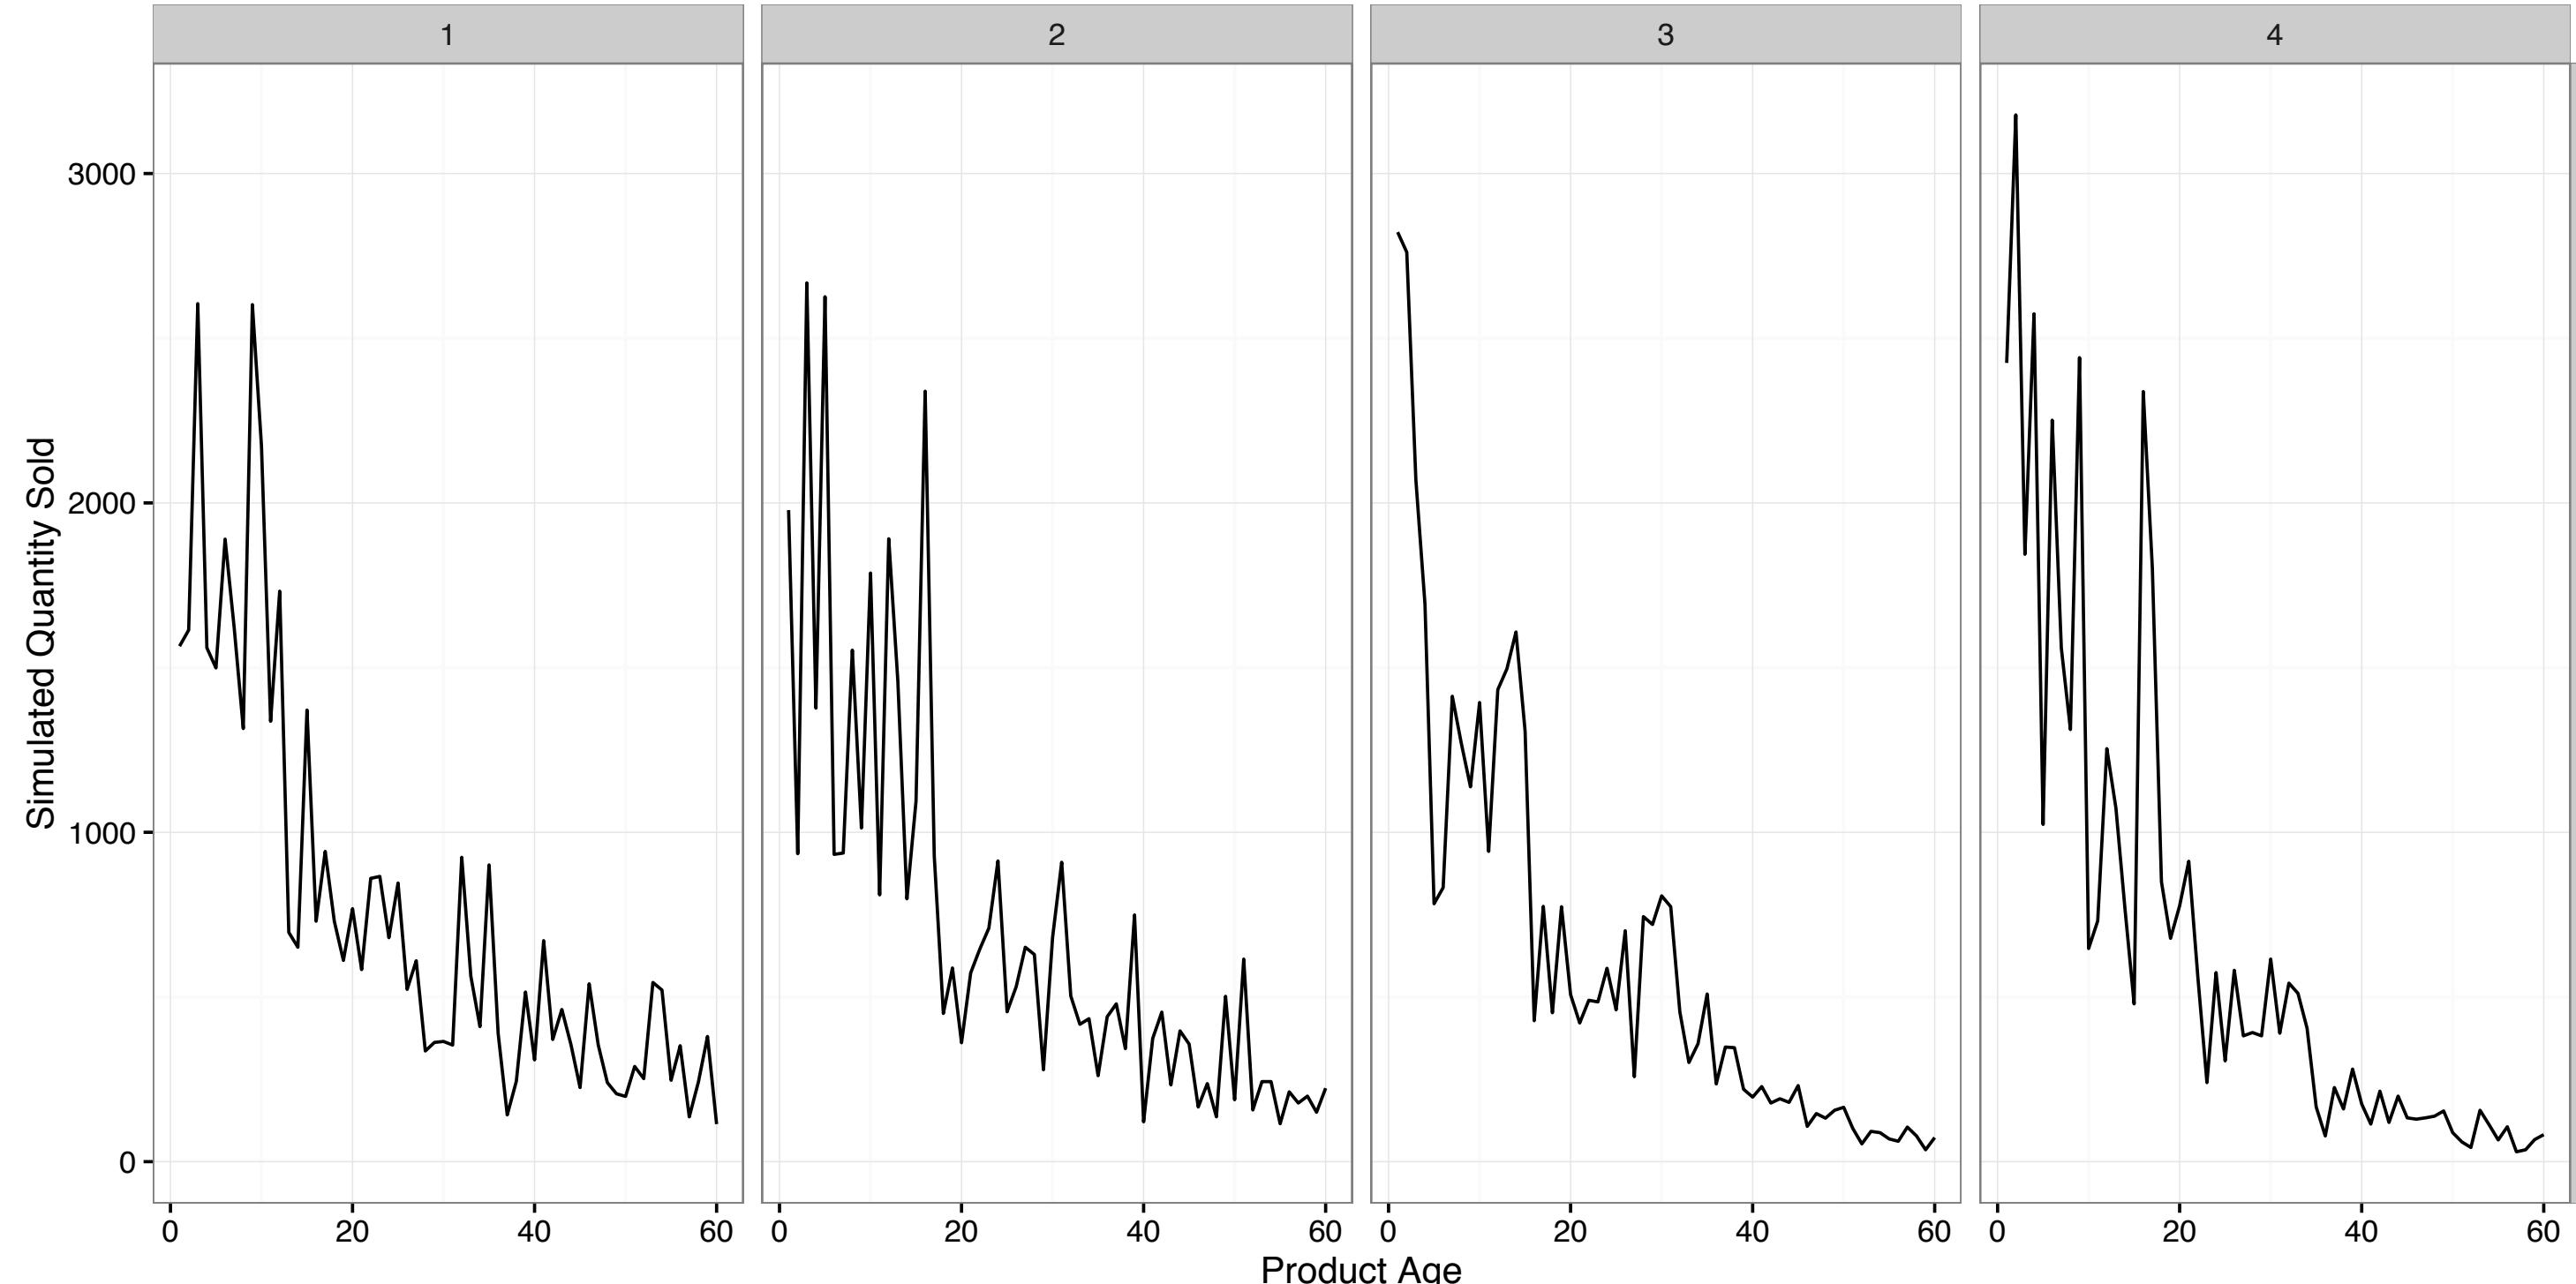
$$\sigma^2 = \mu + \mu^2/\phi$$



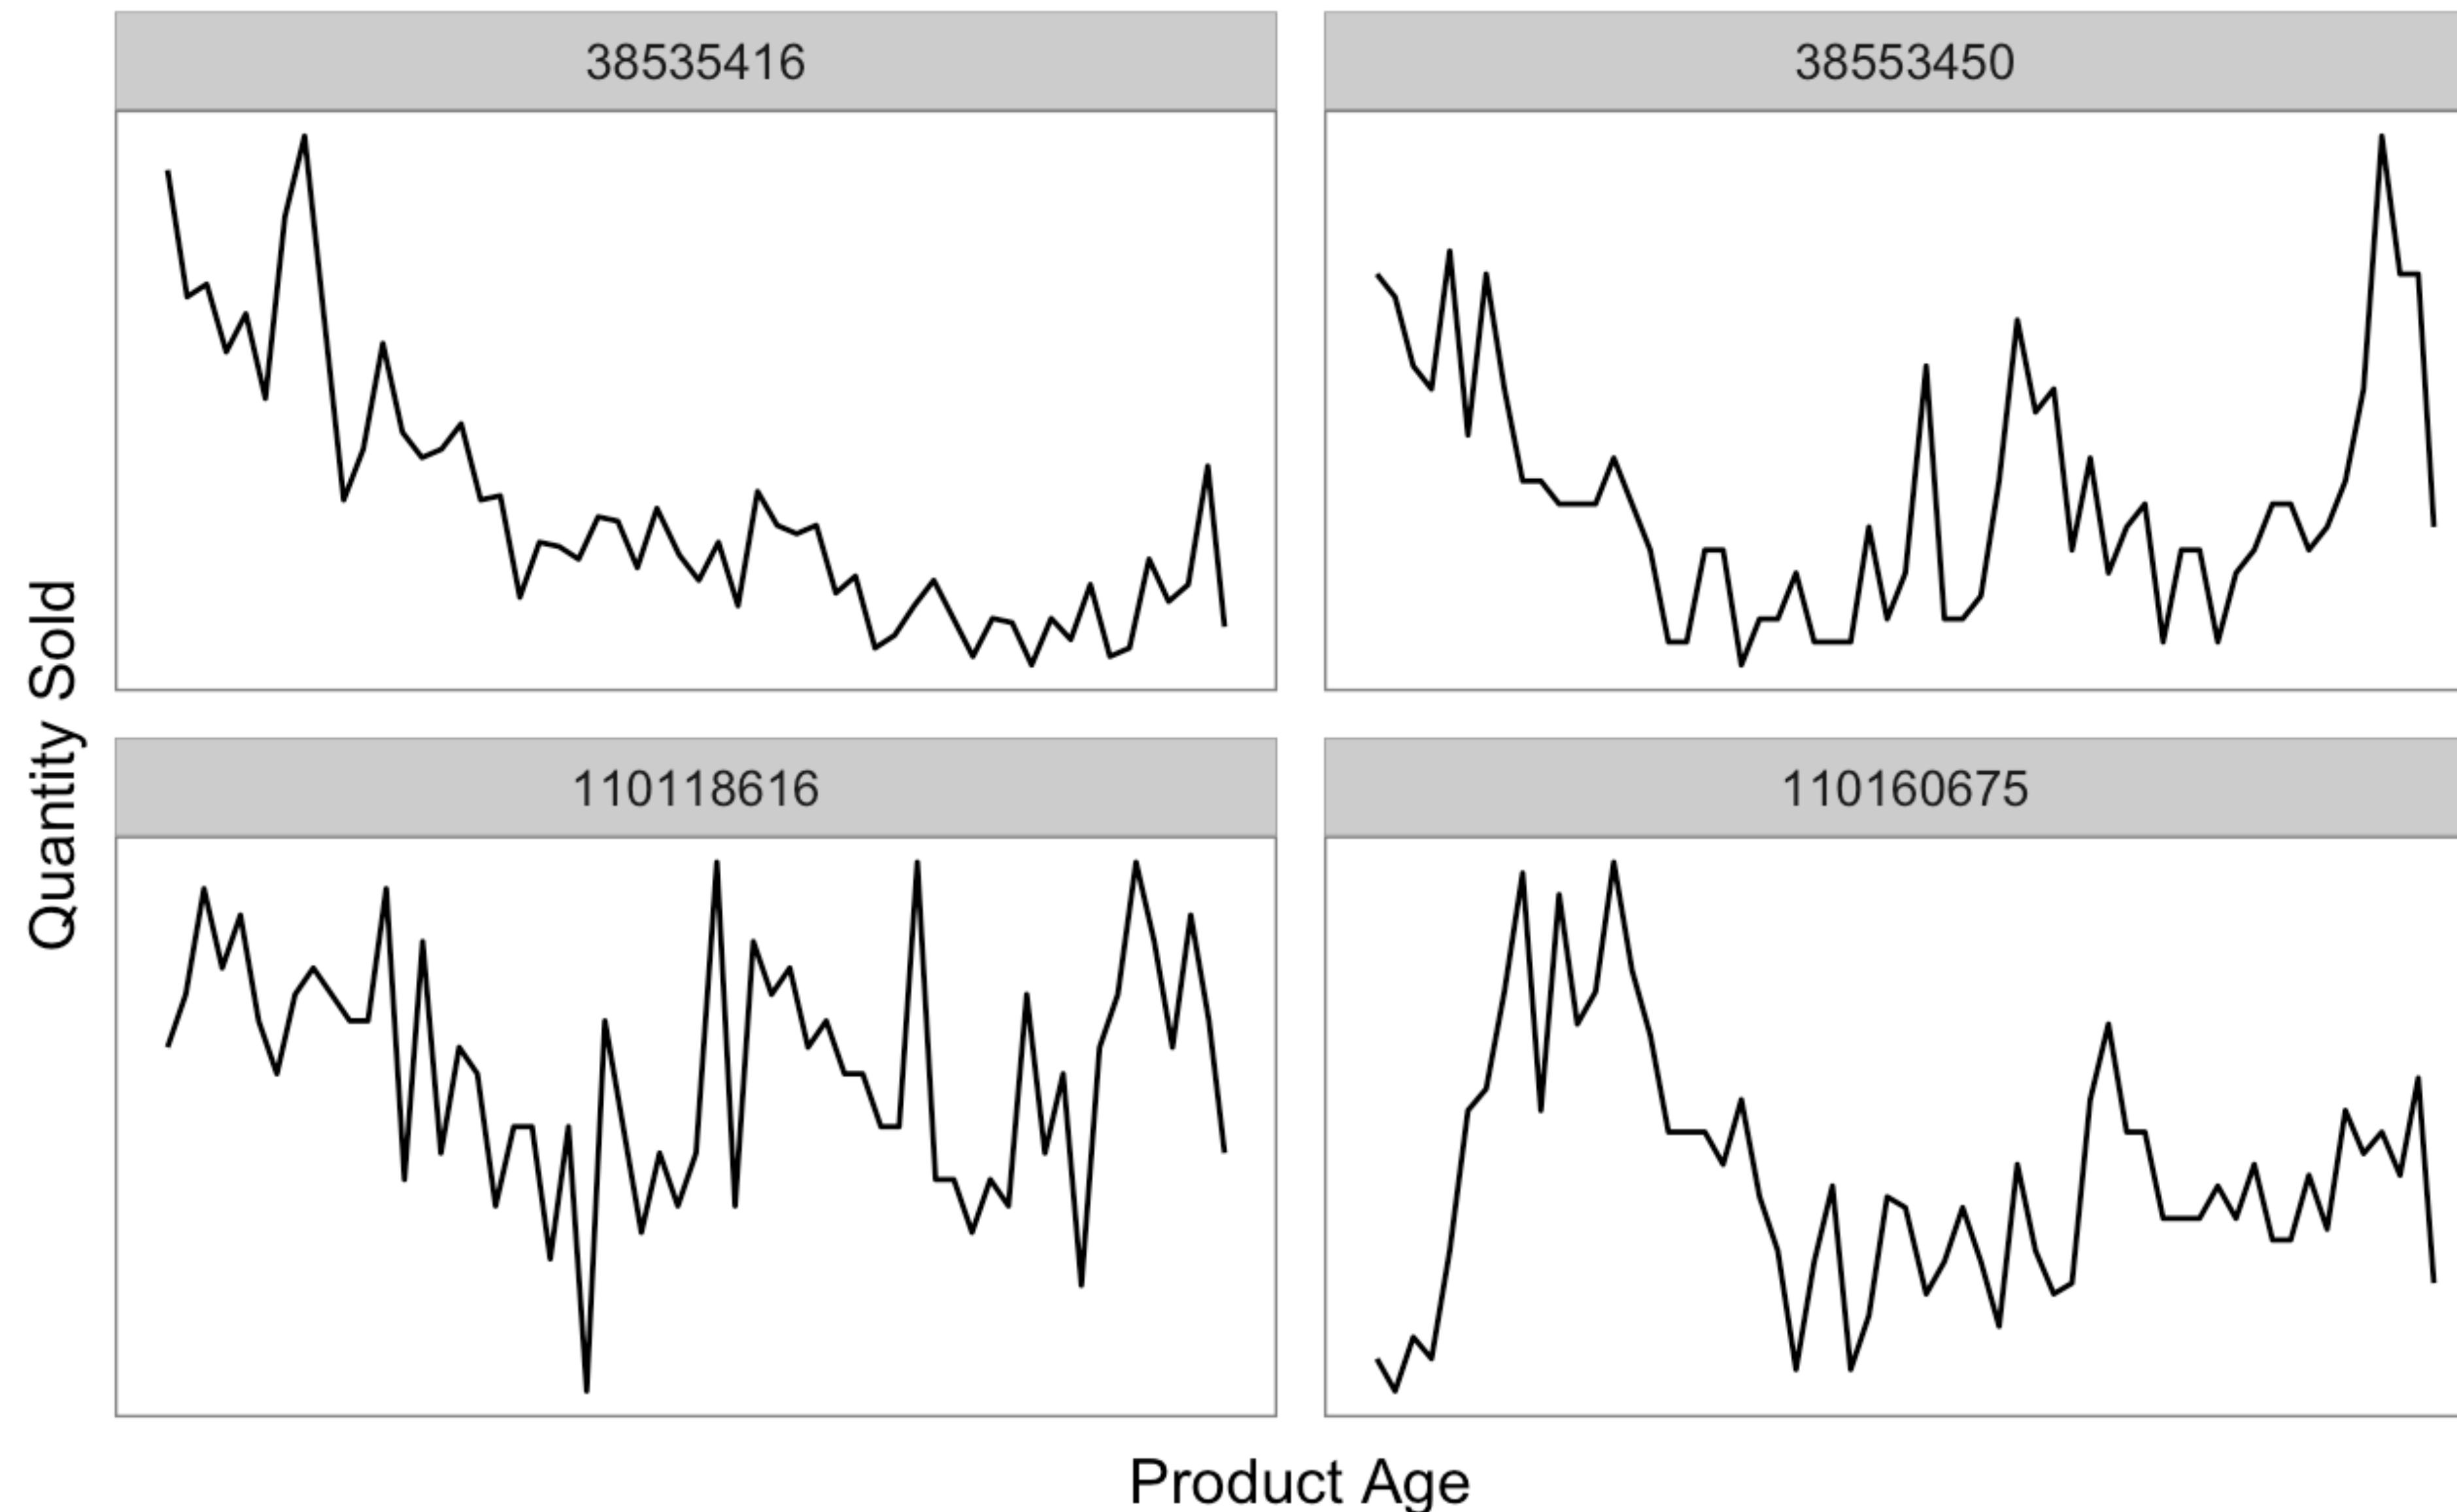
```
simd2 <- hir_data_sim(N_prod = 1,  
global_intercept = 8.5,  
theta = 10,  
qty_process = "negbinom",  
primary_price_process = "none",  
...  
linkinv = exp)
```

# Simulating Data

```
if (process == "normal") {  
  data <- data %>%  
    mutate(qty = linkinv(product_intercept + product_beta_time * days + product_beta_price * price +  
      error_sd * rnorm(sum(n)))) %>%  
    mutate(qty = ifelse(qty <= 0, 0, round(qty)))  
} else { # negative binomial  
  data <- data %>%  
    mutate(mu = linkinv(product_intercept + product_beta_time * day + product_beta_price * price))  
    qty = MASS::rnegbin(n = sum(n), mu = mu, theta = theta))  
}
```



# What About the Real Data?



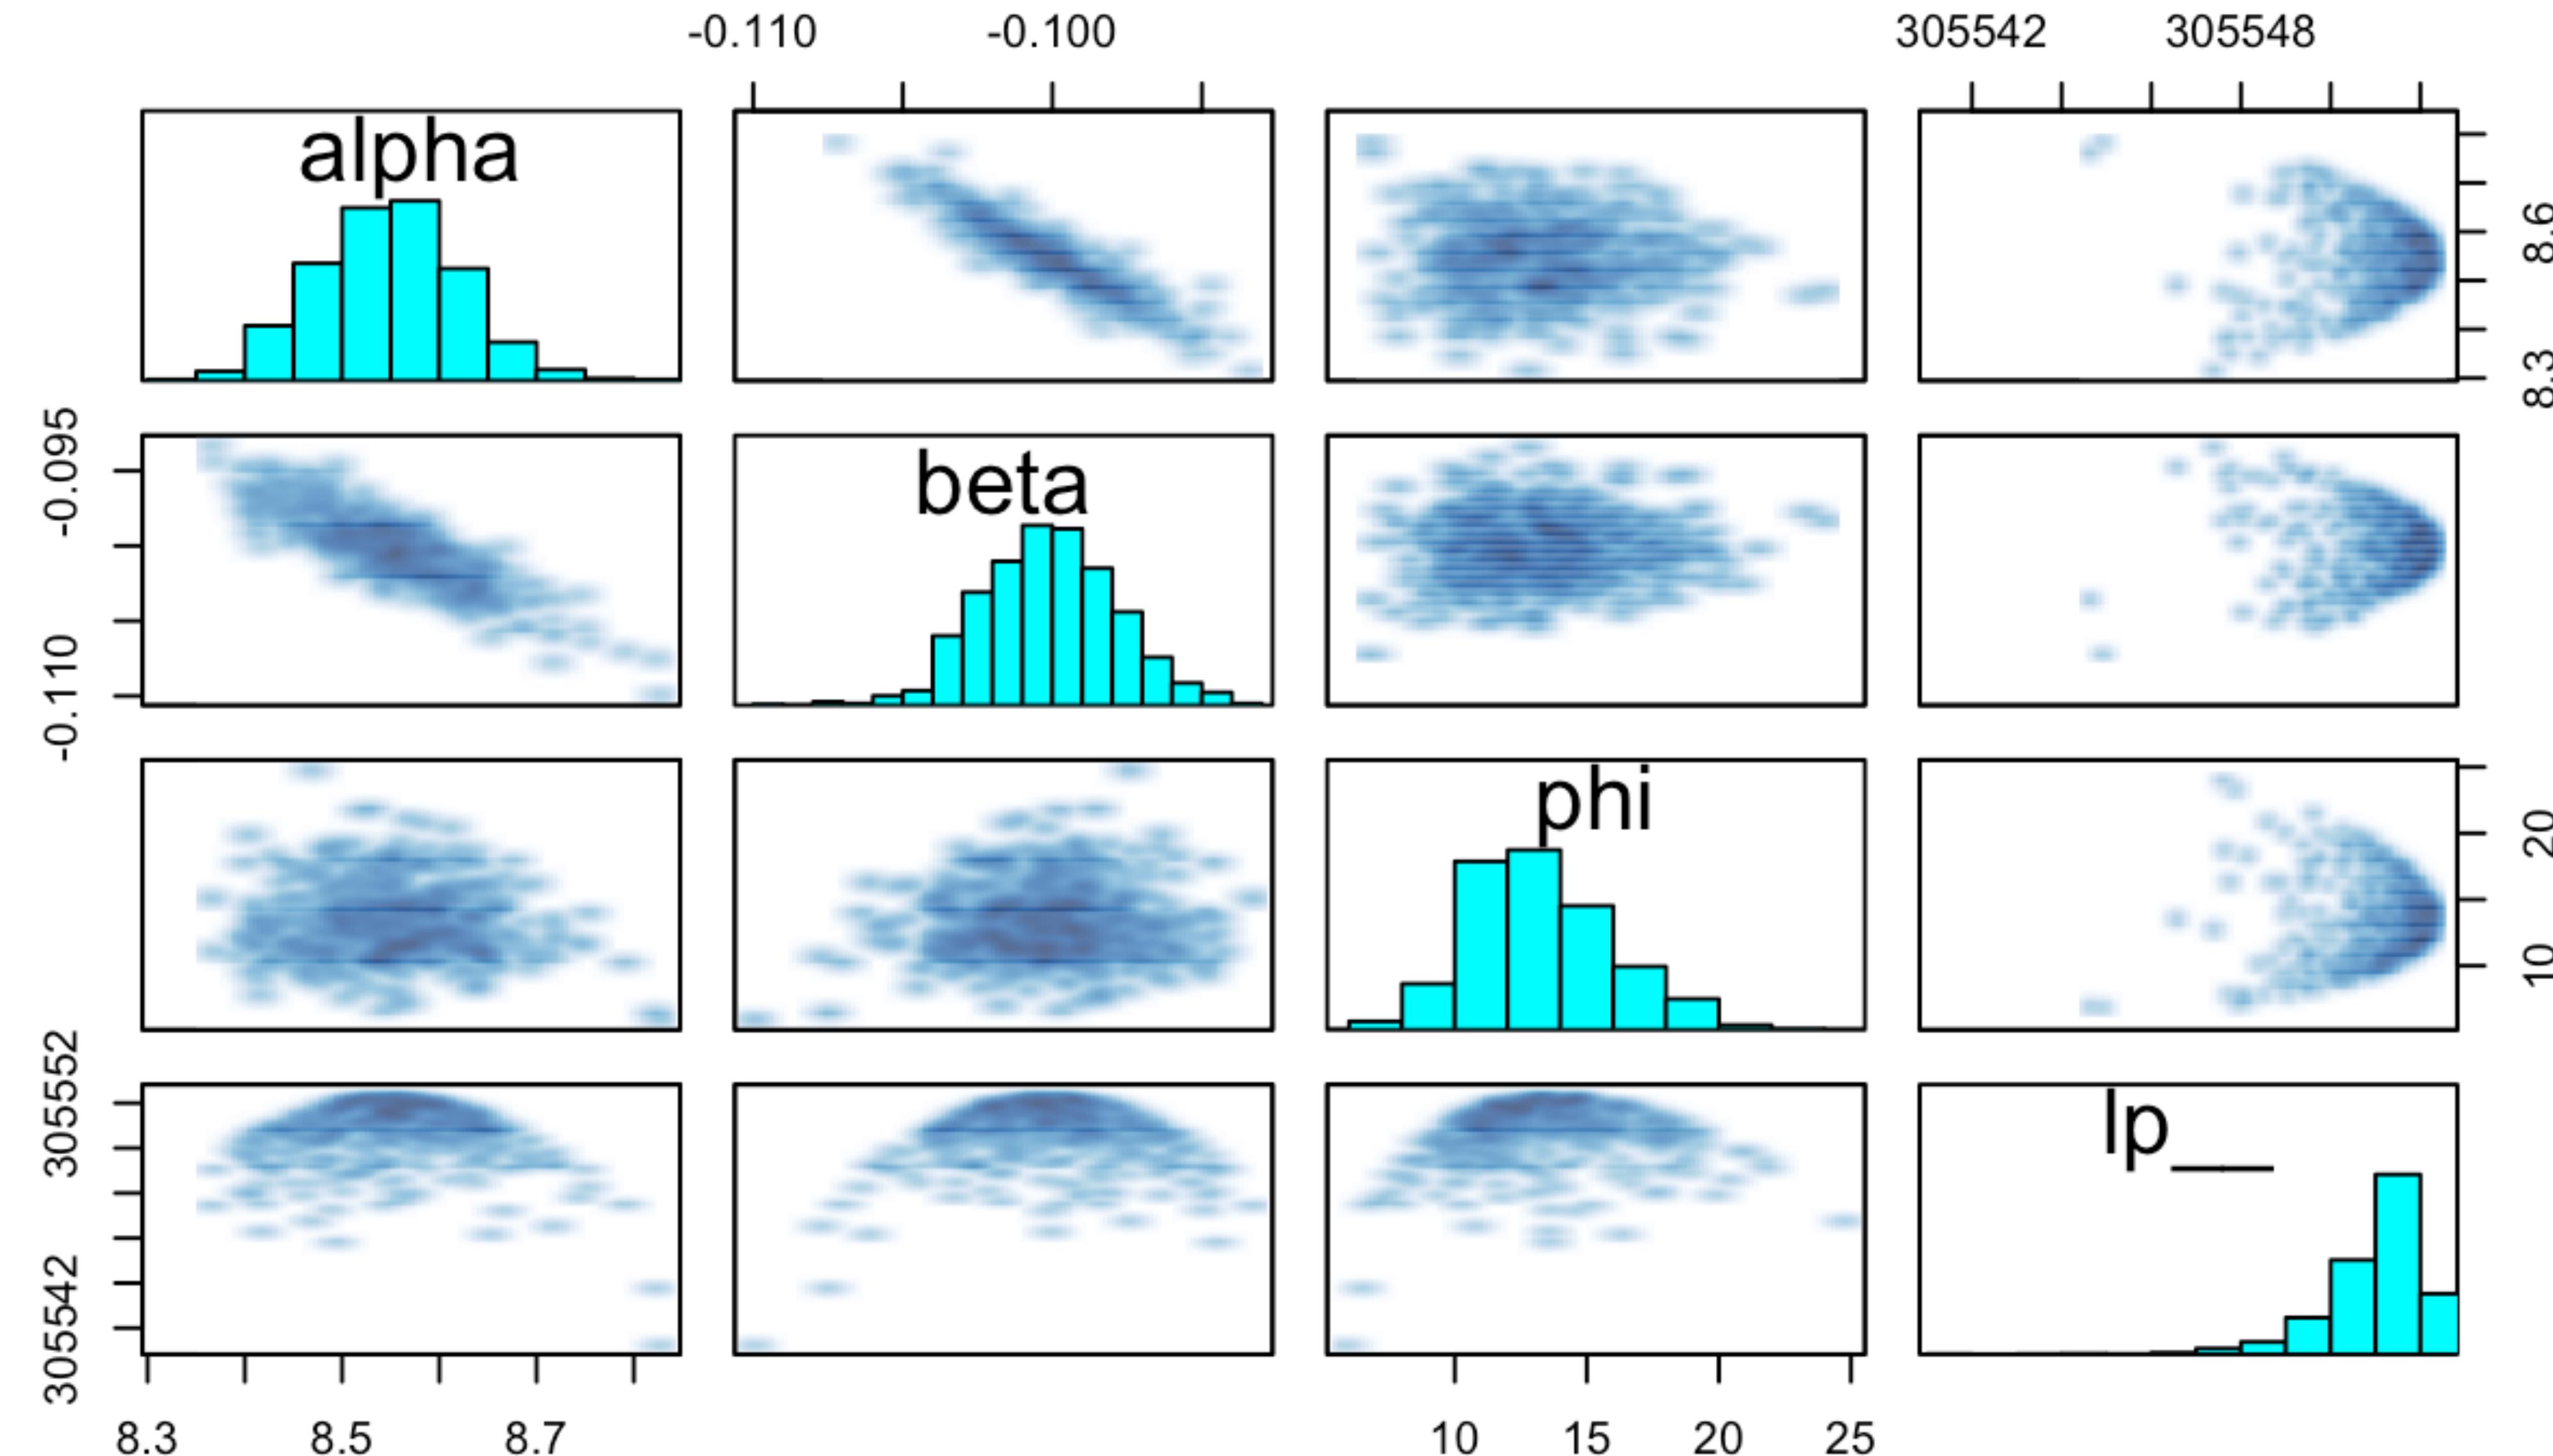
# Baseline Stan Model for Single Product

```
data {  
    int<lower=0> N;  
    int<lower=0> y[N];  
    vector[N] t;  
}  
parameters {  
    real alpha;          // overall mean  
    real beta;           // time beta  
    real<lower=0> phi;  // dispersion  
}  
model {  
    vector[N] eta;  
    // linear predictor  
    eta = alpha + t * beta;  
    // priors  
    alpha ~ normal(0, 10);  
    phi ~ cauchy(0, 2.5);  
    beta ~ normal(0, 1);  
    // likelihood  
    y ~ neg_binomial_2_log(eta, phi);  
}
```

```
simd2_m2 <- stan('m2_self_stan_nbino.stan'  
                   data = list(N = nrow(simd2$data),  
                               y = simd2$data$qty,  
                               t = simd2$data$day),  
                   control = list(stepsize = 0.01,  
                                 adapt_delta = 0.99),  
                   cores = 4,  
                   iter = 400)  
  
# truth: alpha = 8.5, beta = -0.10, phi = 10  
samples <- rstan::extract(simd2_m2,  
                           pars = c('alpha',  
                                   'beta',  
                                   'phi'))  
  
> lapply(samples, quantile)  
$alpha  
  0%  25%  50%  75% 100%  
 8.3  8.4  8.5  8.6  8.8  
  
$beta  
  0%   25%   50%   75% 100%  
-0.107 -0.102 -0.100 -0.099 -0.092  
  
$phi  
  0%  25%  50%  75% 100%  
 6.2 10.1 11.5 13.0 24.1
```

# Looking at Posterior Draws

```
> pairs(simd2_m2)
```



# Diagnostics with ShinyStan

Parameter

alpha

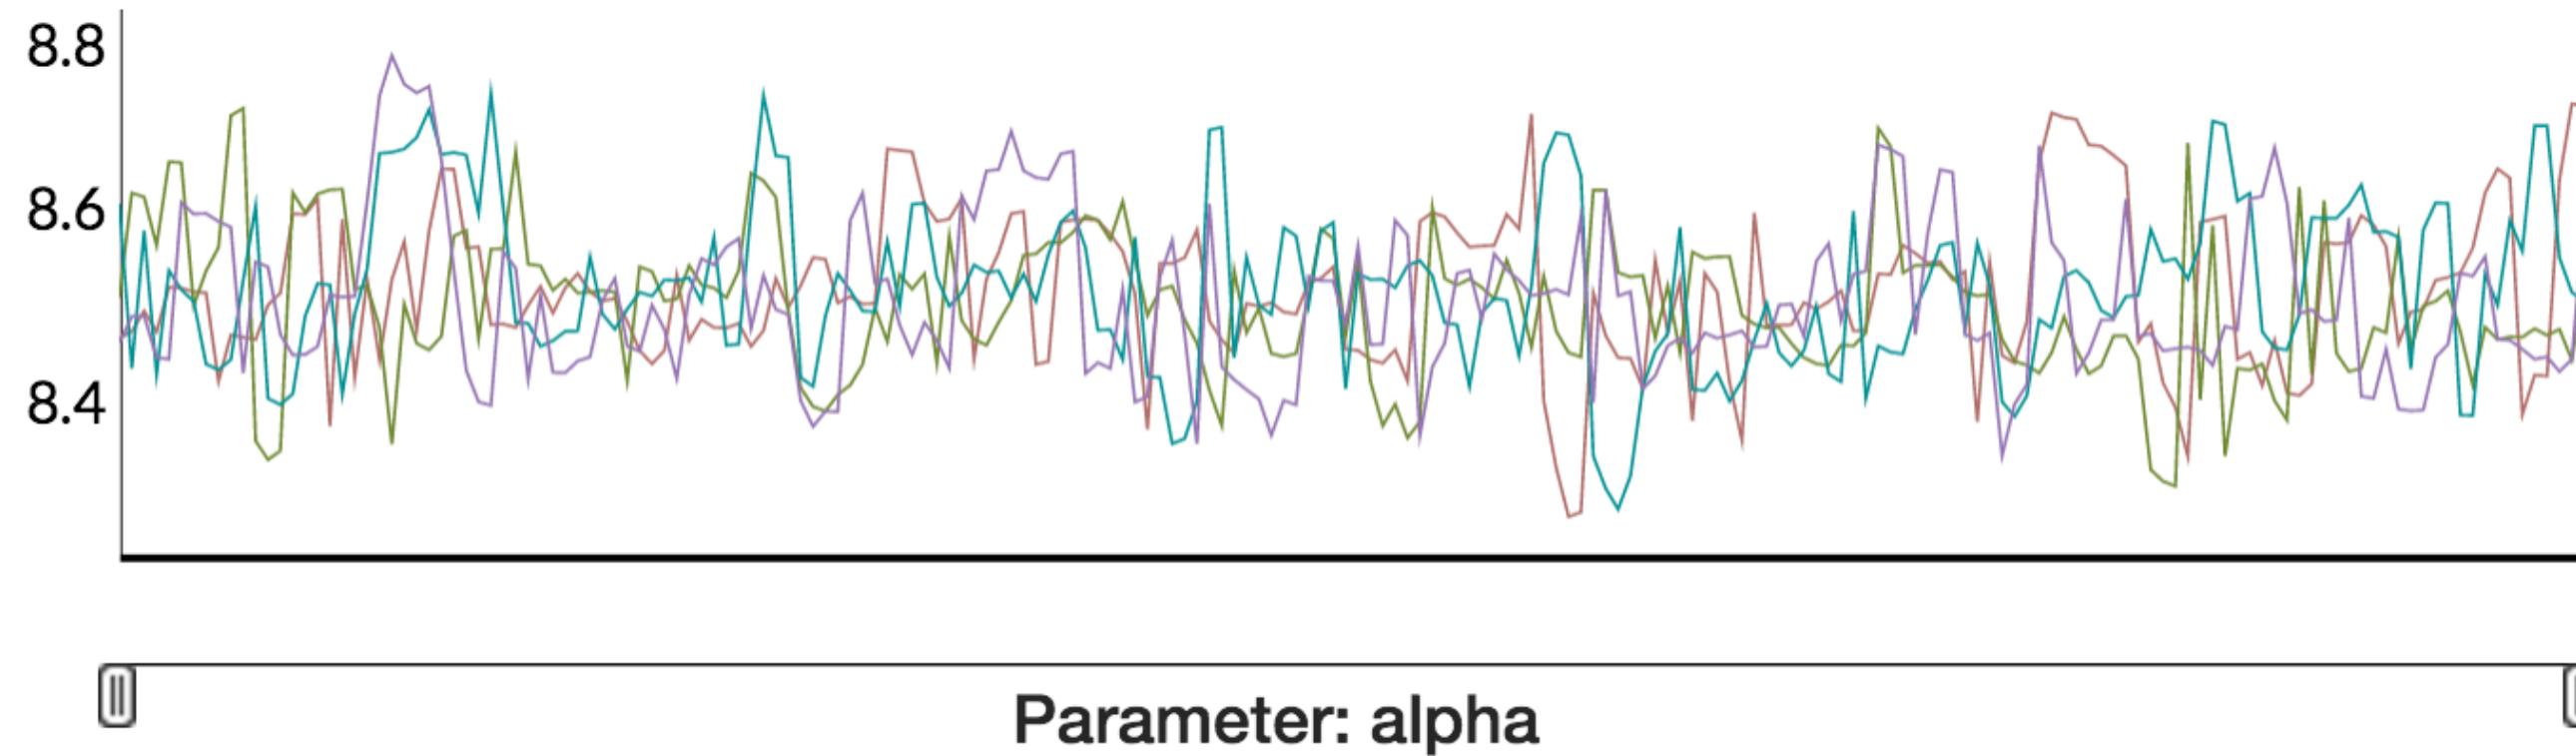
Transformation

identity

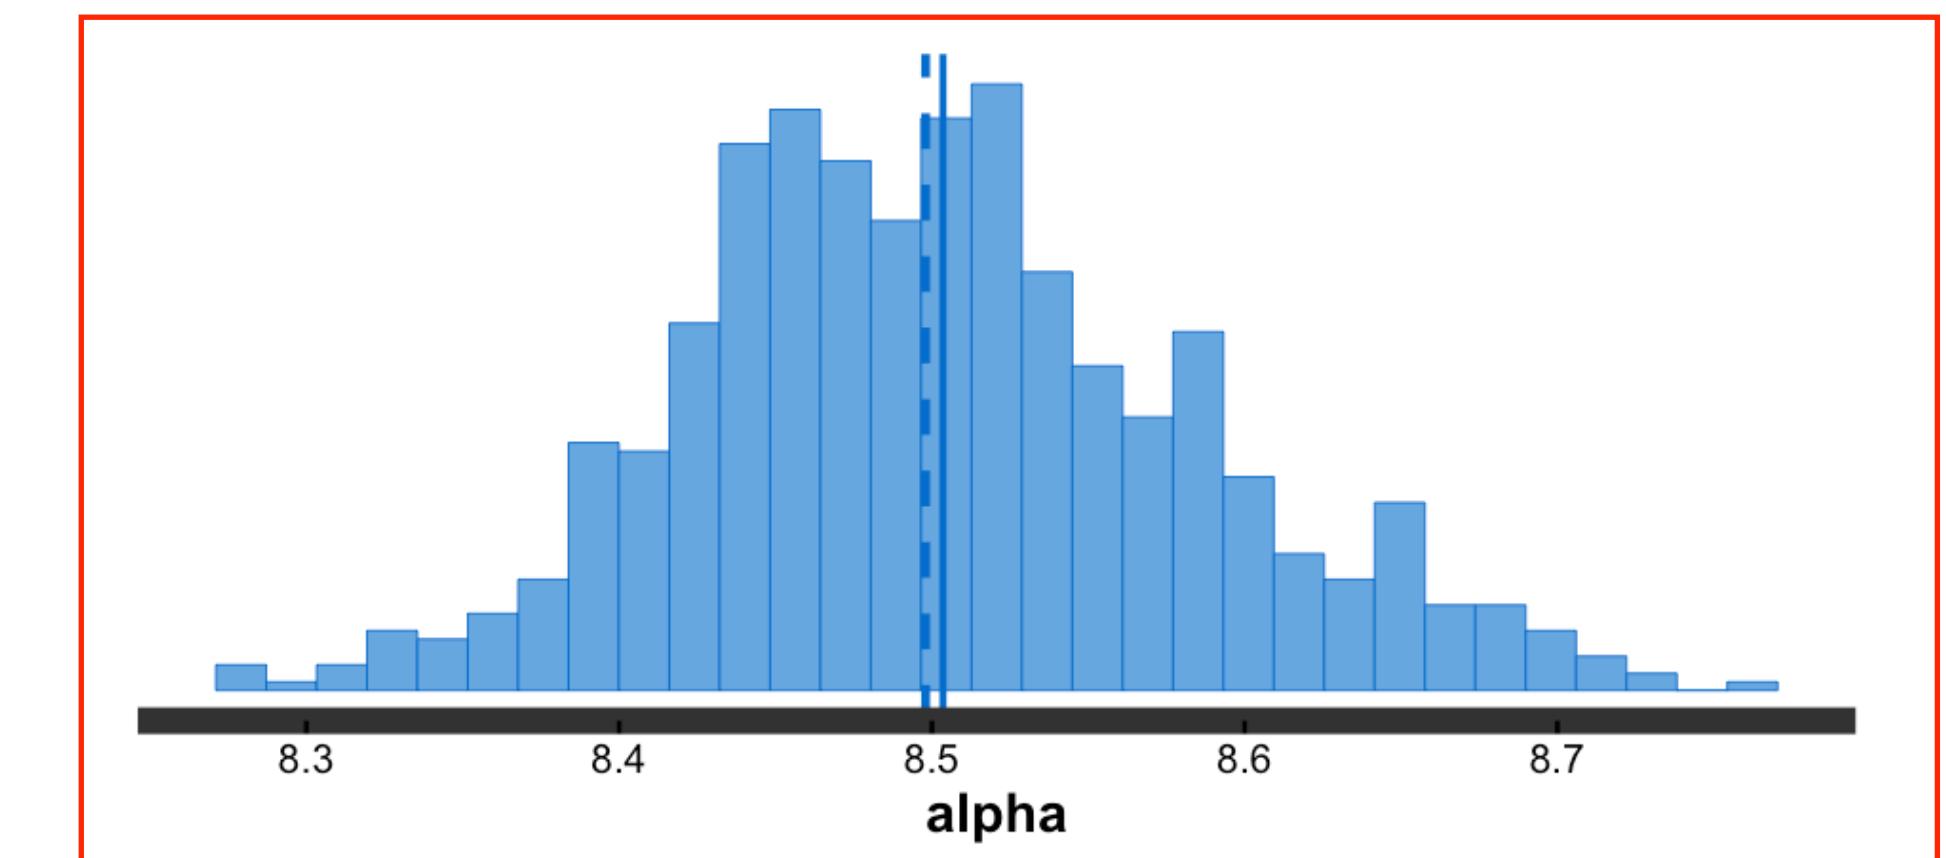
Transform

Use your mouse or the sliders to select areas in the traceplot to zoom into. The other plots on the screen will update accordingly.

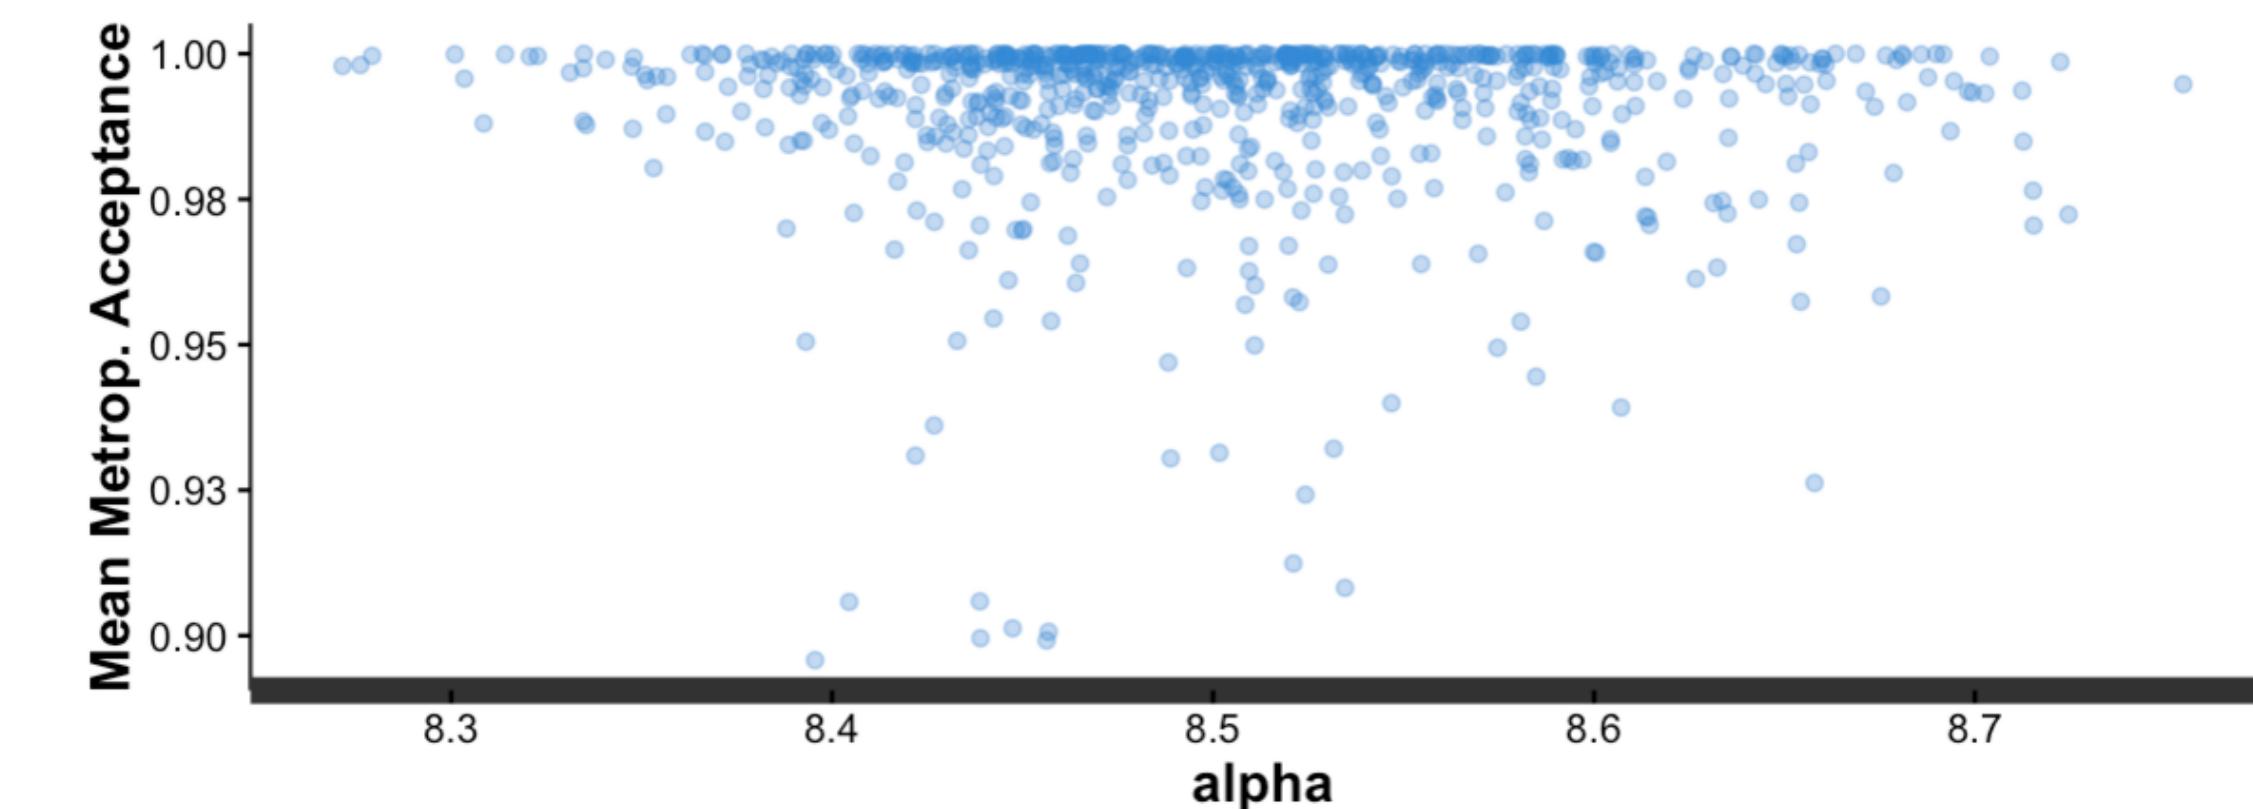
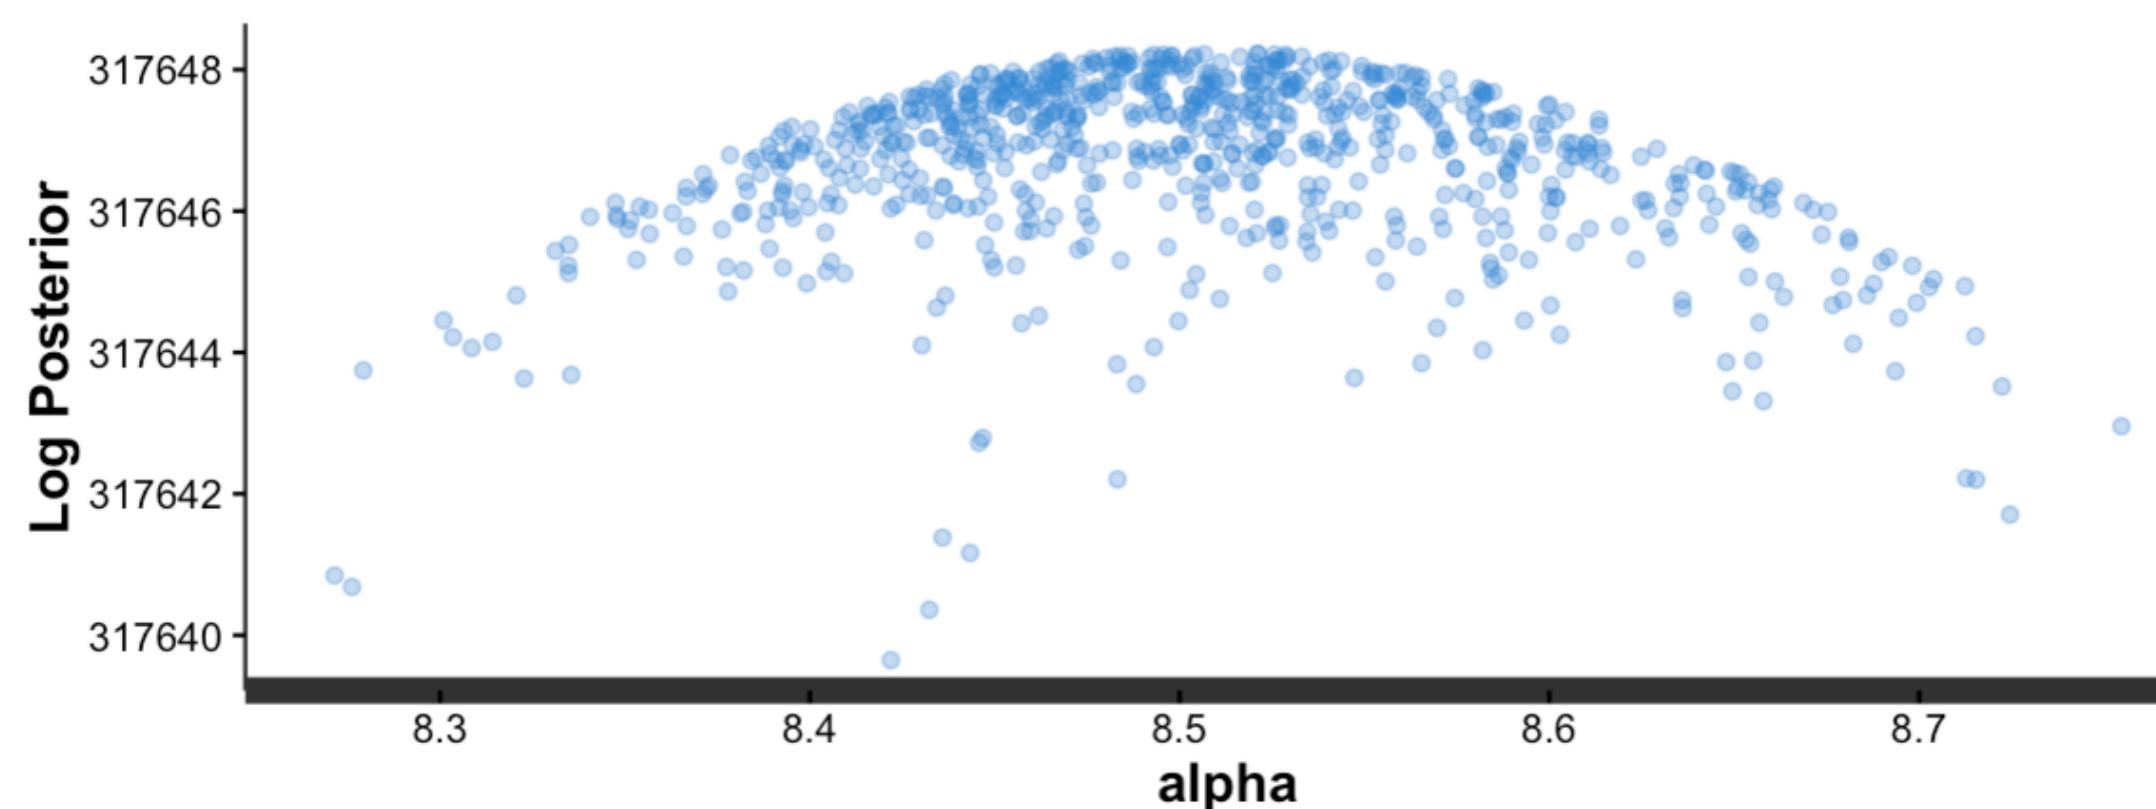
Double-click to reset.



Lines are mean (solid) and median (dashed)



Large red points indicate which (if any) iterations encountered a divergent transition. Yellow indicates a transition hitting the maximum treedepth.

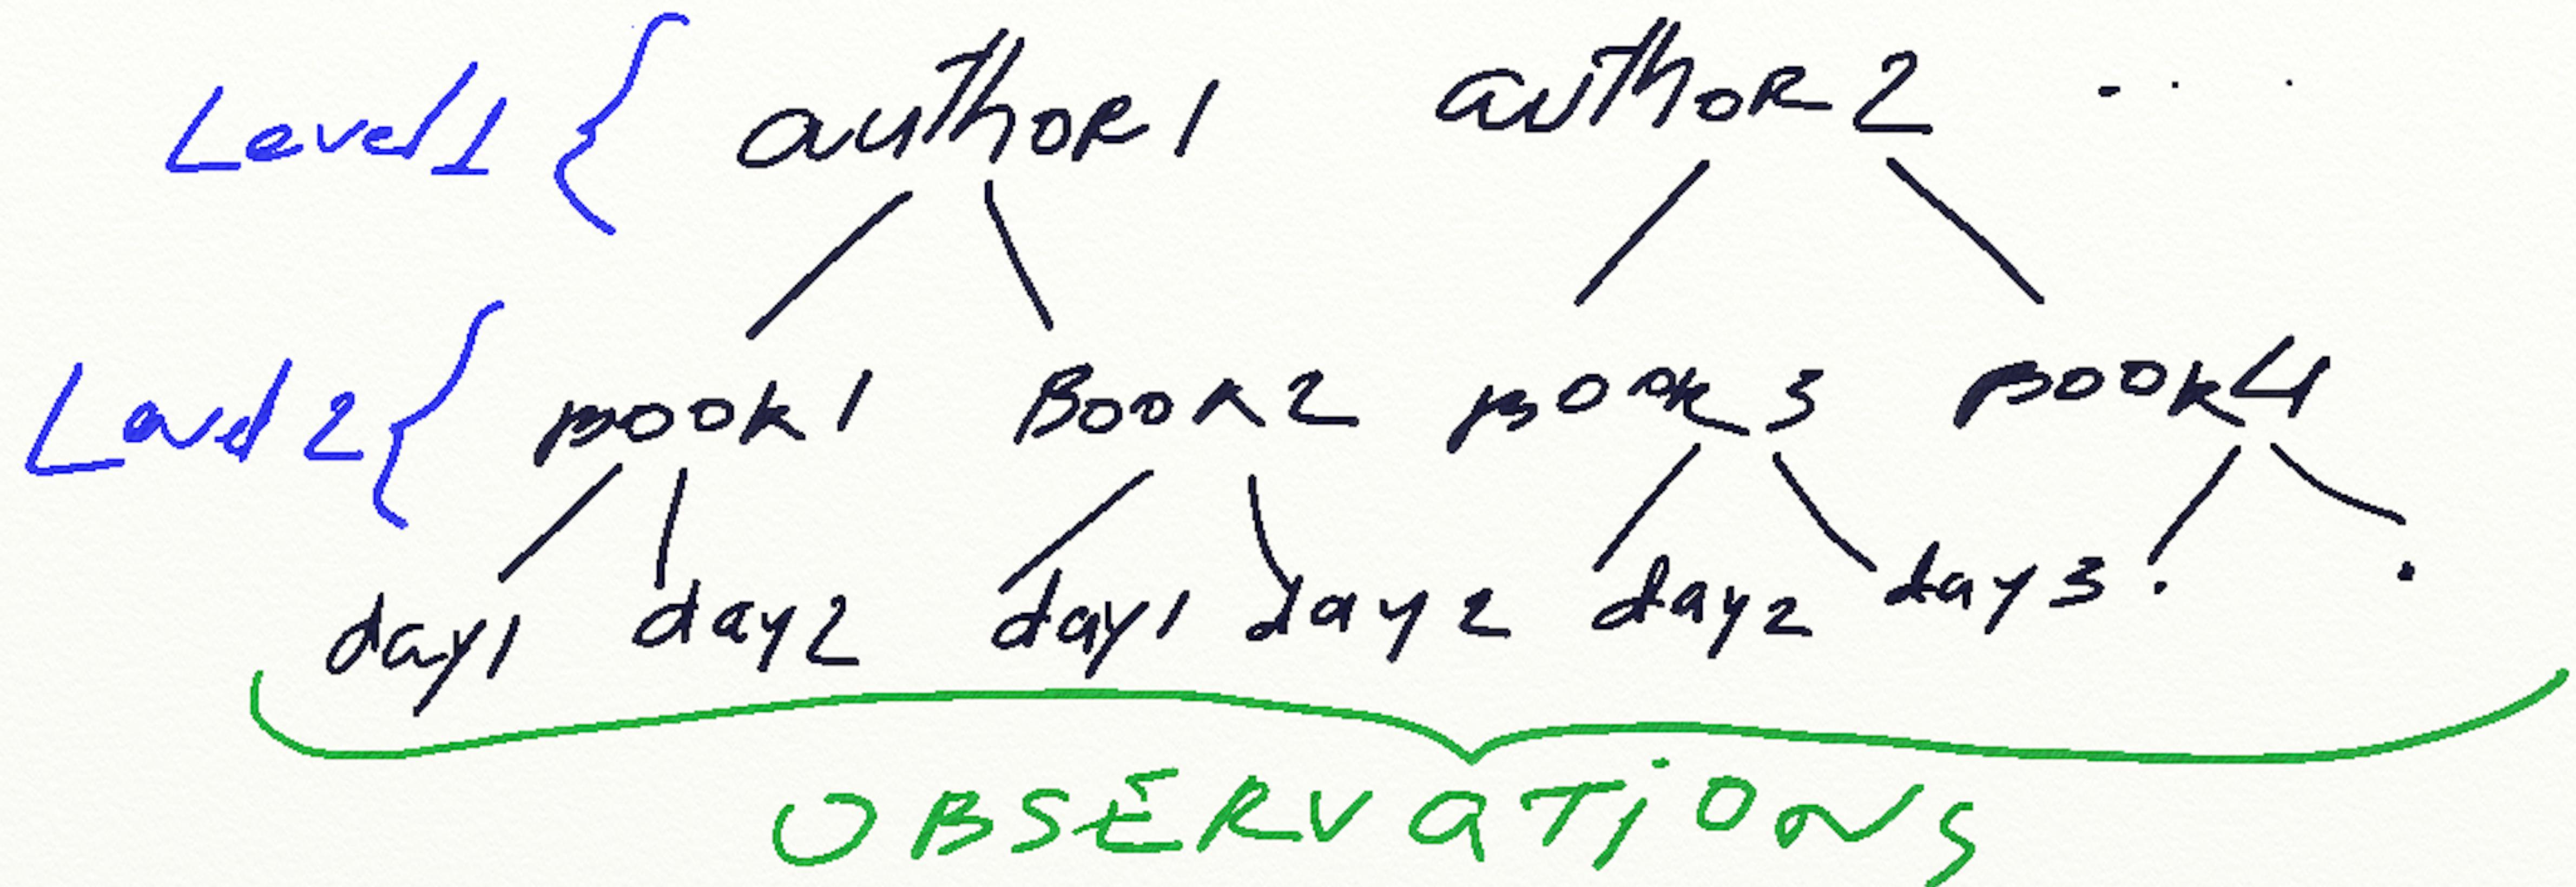


# Product Pricing

---

Introduction to Pooling and Hirarchical Models

# We Have Multiple Products, Authors, Genres



# Hierarchical Pooling in One Slide

$$\hat{\alpha}_j^{multilevel} \approx \frac{\frac{n_j}{\sigma_y^2} \bar{y}_j + \frac{1}{\sigma_\alpha^2} \bar{y}_{all}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}}$$

Number of observations for book j  
Indexes books  
Estimate of average sales for book j  
Average sales for book j  
Within-book variance  
Variance among average sales of different books  
Average sales across all books

The diagram illustrates the hierarchical pooling formula. It shows the components of the multilevel estimate  $\hat{\alpha}_j^{multilevel}$  and their relationships to the data and parameters. The formula is:

$$\hat{\alpha}_j^{multilevel} \approx \frac{\frac{n_j}{\sigma_y^2} \bar{y}_j + \frac{1}{\sigma_\alpha^2} \bar{y}_{all}}{\frac{n_j}{\sigma_y^2} + \frac{1}{\sigma_\alpha^2}}$$

The components are labeled as follows:

- Number of observations for book j (points to  $n_j/\sigma_y^2$ )
- Average sales for book j (points to  $\bar{y}_j$ )
- Within-book variance (points to  $n_j/\sigma_y^2$ )
- Variance among average sales of different books (points to  $1/\sigma_\alpha^2$ )
- Average sales across all books (points to  $\bar{y}_{all}$ )

# Multi-Level Models using Lmer Syntax

| Formula                             | Alternative                           | Meaning   |
|-------------------------------------|---------------------------------------|---|
| $(1 \mid g)$                        | $1 + (1 \mid g)$                      | Random intercept<br>with fixed mean               |
| $0 + \text{offset}(o) + (1 \mid g)$ | $-1 + \text{offset}(o) + (1 \mid g)$  | Random intercept<br>with <i>a priori</i> means    |
| $(1 \mid g1/g2)$                    | $(1 \mid g1) + (1 \mid g1:g2)$        | Intercept varying<br>among g1 and g2<br>within g1 |
| $(1 \mid g1) + (1 \mid g2)$         | $1 + (1 \mid g1) + (1 \mid g2)$       | Intercept varying<br>among g1 and g2              |
| $x + (x \mid g)$                    | $1 + x + (1 + x \mid g)$              | Correlated random<br>intercept and slope          |
| $x + (x \parallel g)$               | $1 + x + (1 \mid g) + (0 + x \mid g)$ | Uncorrelated random<br>intercept and slope        |

Table 2: Examples of the right-hand sides of mixed-effects model formulas. The names of grouping factors are denoted  $g$ ,  $g1$ , and  $g2$ , and covariates and *a priori* known offsets as  $x$  and  $o$ .

# Fitting Multi-Level Models in rstanarm

```
fit <- stan_glmer.nb(qty ~ product_age + price + price_sqr +  
                      (1 + product_age + price + price_sqr | product),  
                      algorithm = "sampling",  
                      seed = 123,  
                      cores = 4,  
                      iter = 600,  
                      data = data)
```

“Fixed Effects”

Varying Intercepts

Varying Slopes

Fit using MCMC

Random Seed

1 Core per Chain

Number of Iterations per Chain

The diagram illustrates the components of the R code for fitting a multi-level model. It uses green arrows to point from descriptive labels to specific parts of the code. The labels are: “Fixed Effects”, Varying Intercepts, Varying Slopes, Fit using MCMC, Random Seed, 1 Core per Chain, and Number of Iterations per Chain.

# Prediction and Checking: Posterior Predictive Distribution

- ▶ How can we tell if our model is sufficient for our task?
- ▶ We can simulate from the model and compare to observed data
- ▶ We can predict across interesting co-variates (e.g. change prices and observe how the model predicts qty over time)

$$p(\tilde{y}|y) = \int_{\Theta} p(\tilde{y}|\theta)p(\theta|y)d\theta$$

Posterior Predictive Distribution →  $p(\tilde{y}|y)$

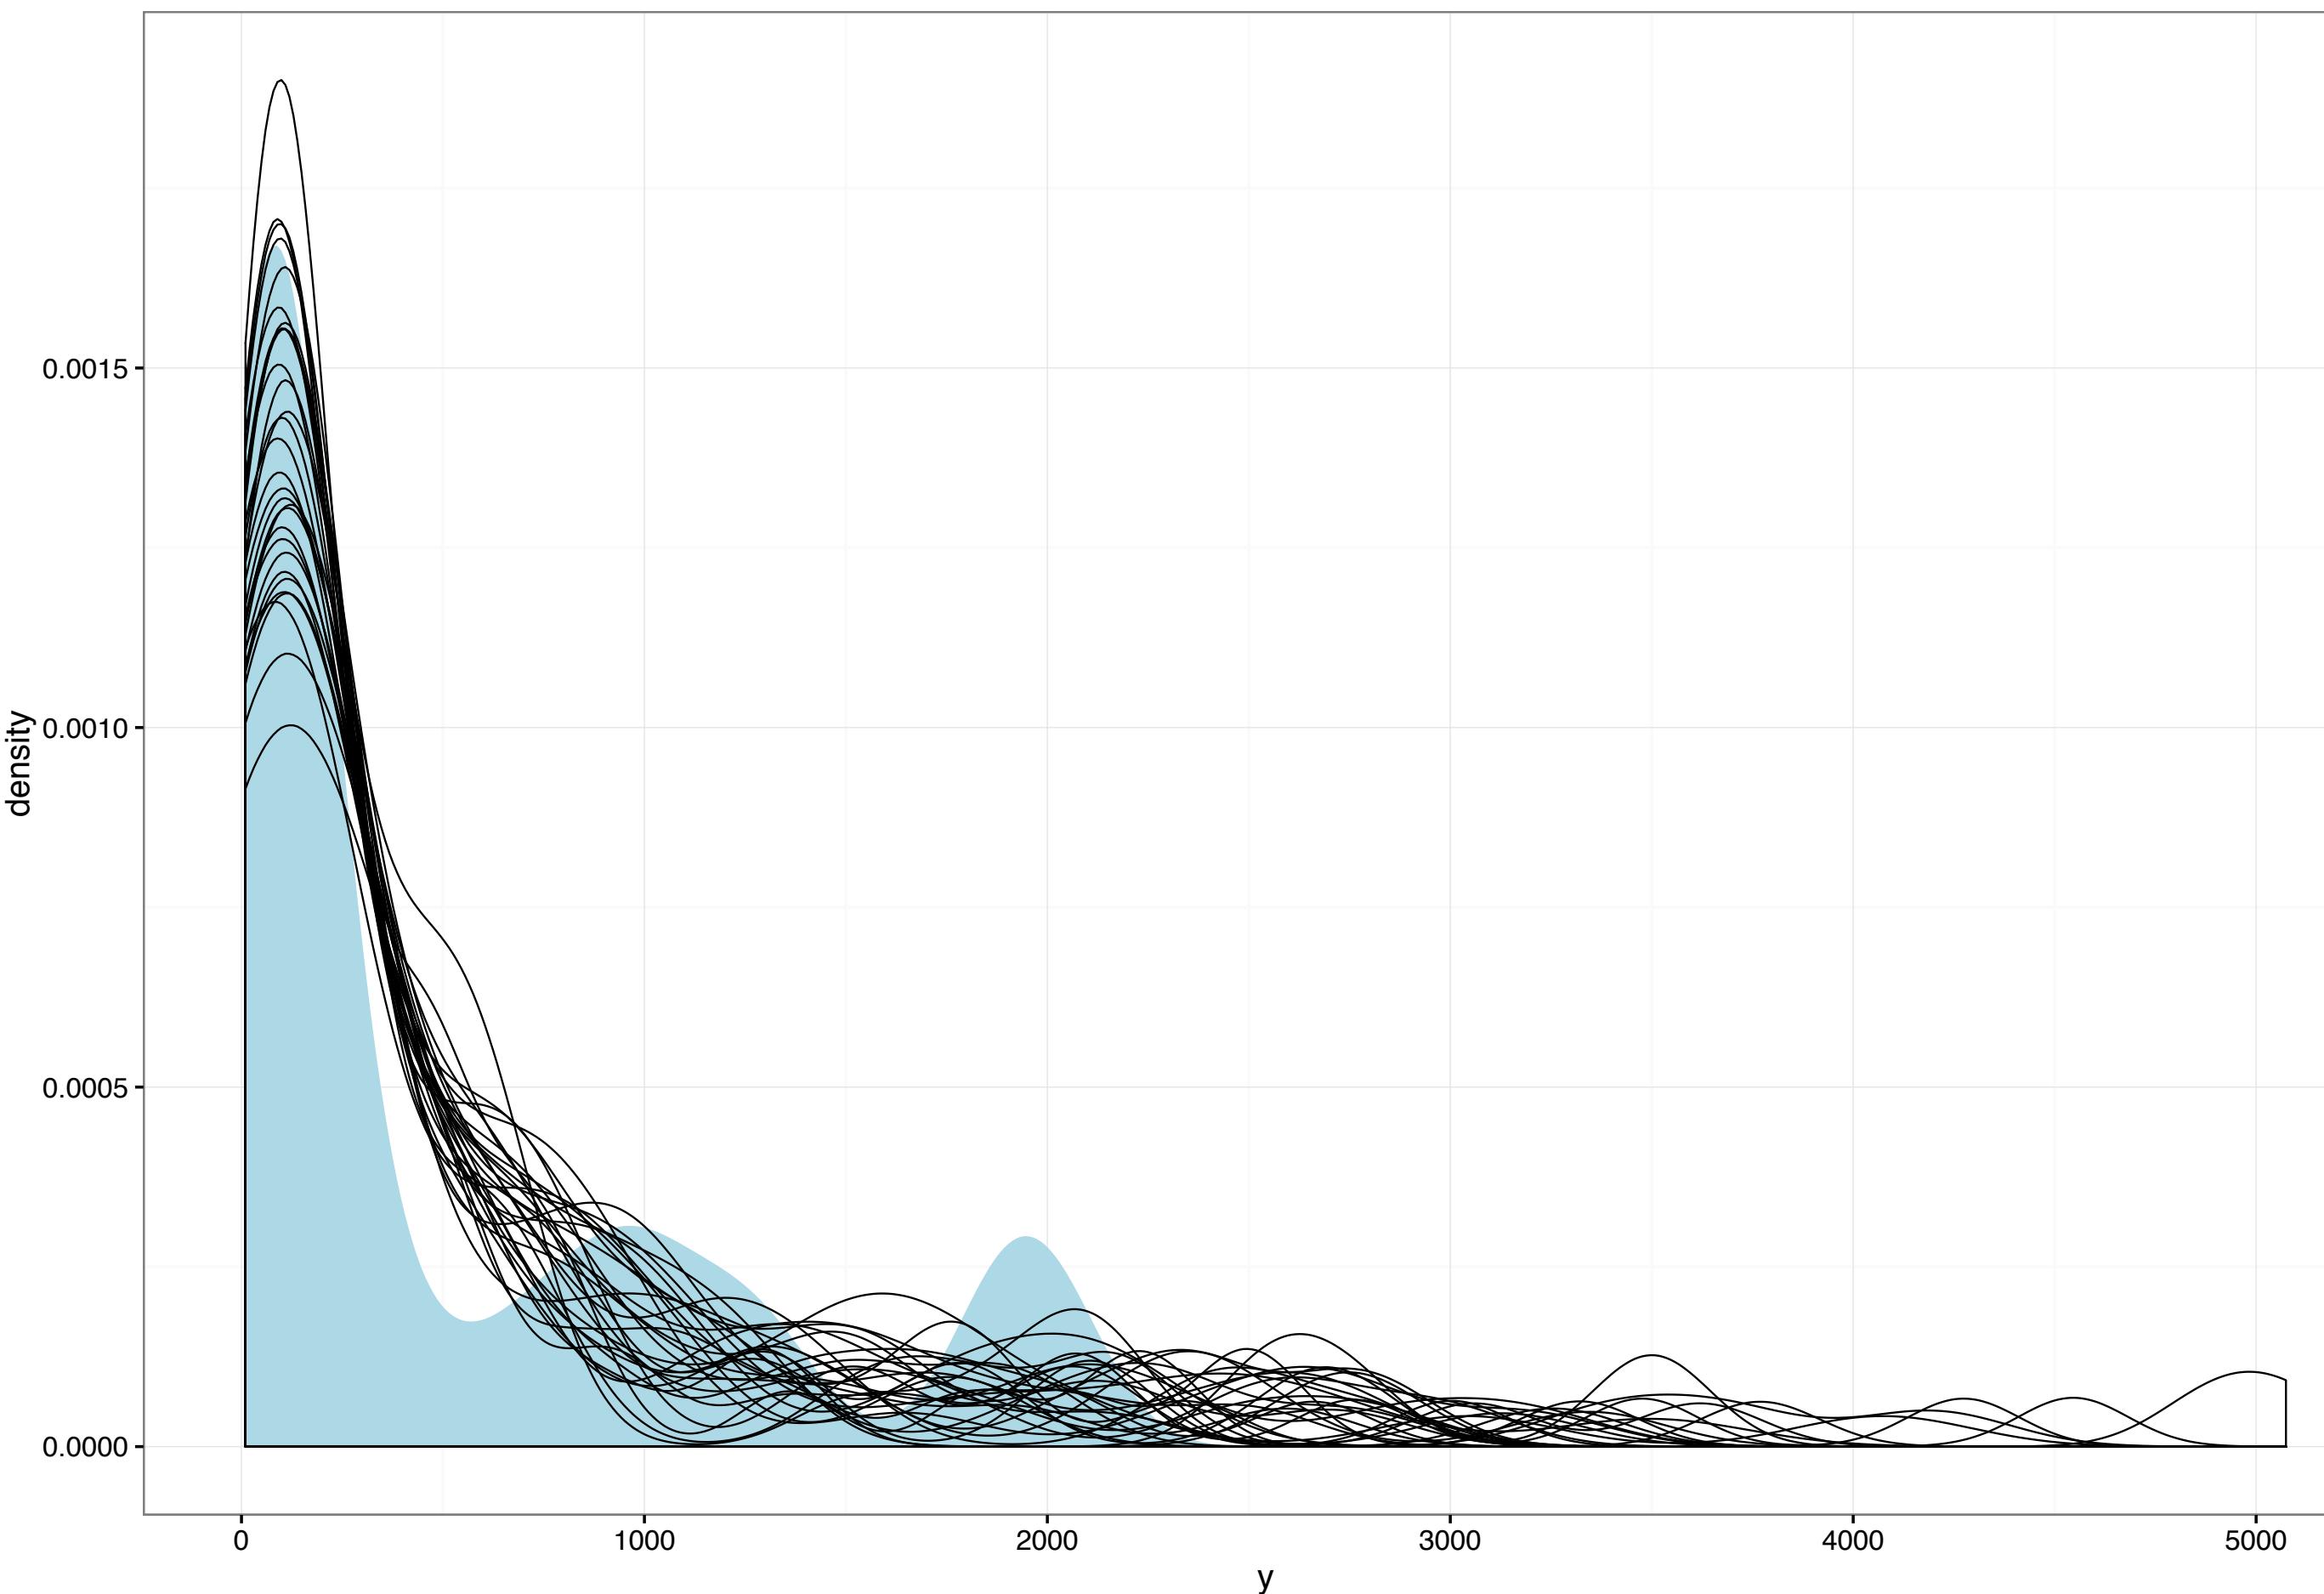
New Data → Data Used to Fit the Model → Likelihood Function → Weighted by the Posterior

Average Over Theta

The diagram illustrates the components of the posterior predictive distribution formula. It shows the flow from 'Posterior Predictive Distribution' through the formula to the individual components: 'New Data', 'Data Used to Fit the Model', 'Likelihood Function', 'Weighted by the Posterior', and finally 'Average Over Theta'.

# Assessing Model Performance: Posterior Predictive Checks, Calibration

```
> pp_check(fit, check = "dist", overlay = TRUE)
```



```
> check_calib(d)
```

|   | in_90     | in_50     |
|---|-----------|-----------|
| 1 | 0.9573893 | 0.7125305 |

```
> check_calib(d, TRUE)
```

Source: local data frame [203 x 3]

|  | id    | in_90 | in_50 |
|--|-------|-------|-------|
|  | (dbl) | (dbl) | (dbl) |

|     |            |           |           |
|-----|------------|-----------|-----------|
| 1   | aaaaaaaa1  | 0.9333333 | 0.7166667 |
| 2   | aaaaaaaa2  | 0.9500000 | 0.8333333 |
| 3   | aaaaaaaa3  | 0.9833333 | 0.8500000 |
| 4   | aaaaaaaa4  | 0.9666667 | 0.6500000 |
| 5   | aaaaaaaa5  | 0.9666667 | 0.7000000 |
| 6   | aaaaaaaa6  | 0.9833333 | 0.8833333 |
| 7   | aaaaaaaa7  | 0.9666667 | 0.6833333 |
| 8   | aaaaaaaa8  | 1.0000000 | 0.7666667 |
| 9   | aaaaaaaa9  | 0.8833333 | 0.6166667 |
| 10  | aaaaaaaa10 | 0.9500000 | 0.8500000 |
| ... | ...        | ...       | ...       |

# Prediction for Observed Prices

In Sample Predictions for 25 Random Products



# Revenue Optimisation

---

Generating Model Counterfactuals

# Generating New Prices

---

```
new_data <- generate_new_prices(data, price_grid = seq(1.99, 14.99, by = 1))
```

```
> new_data
```

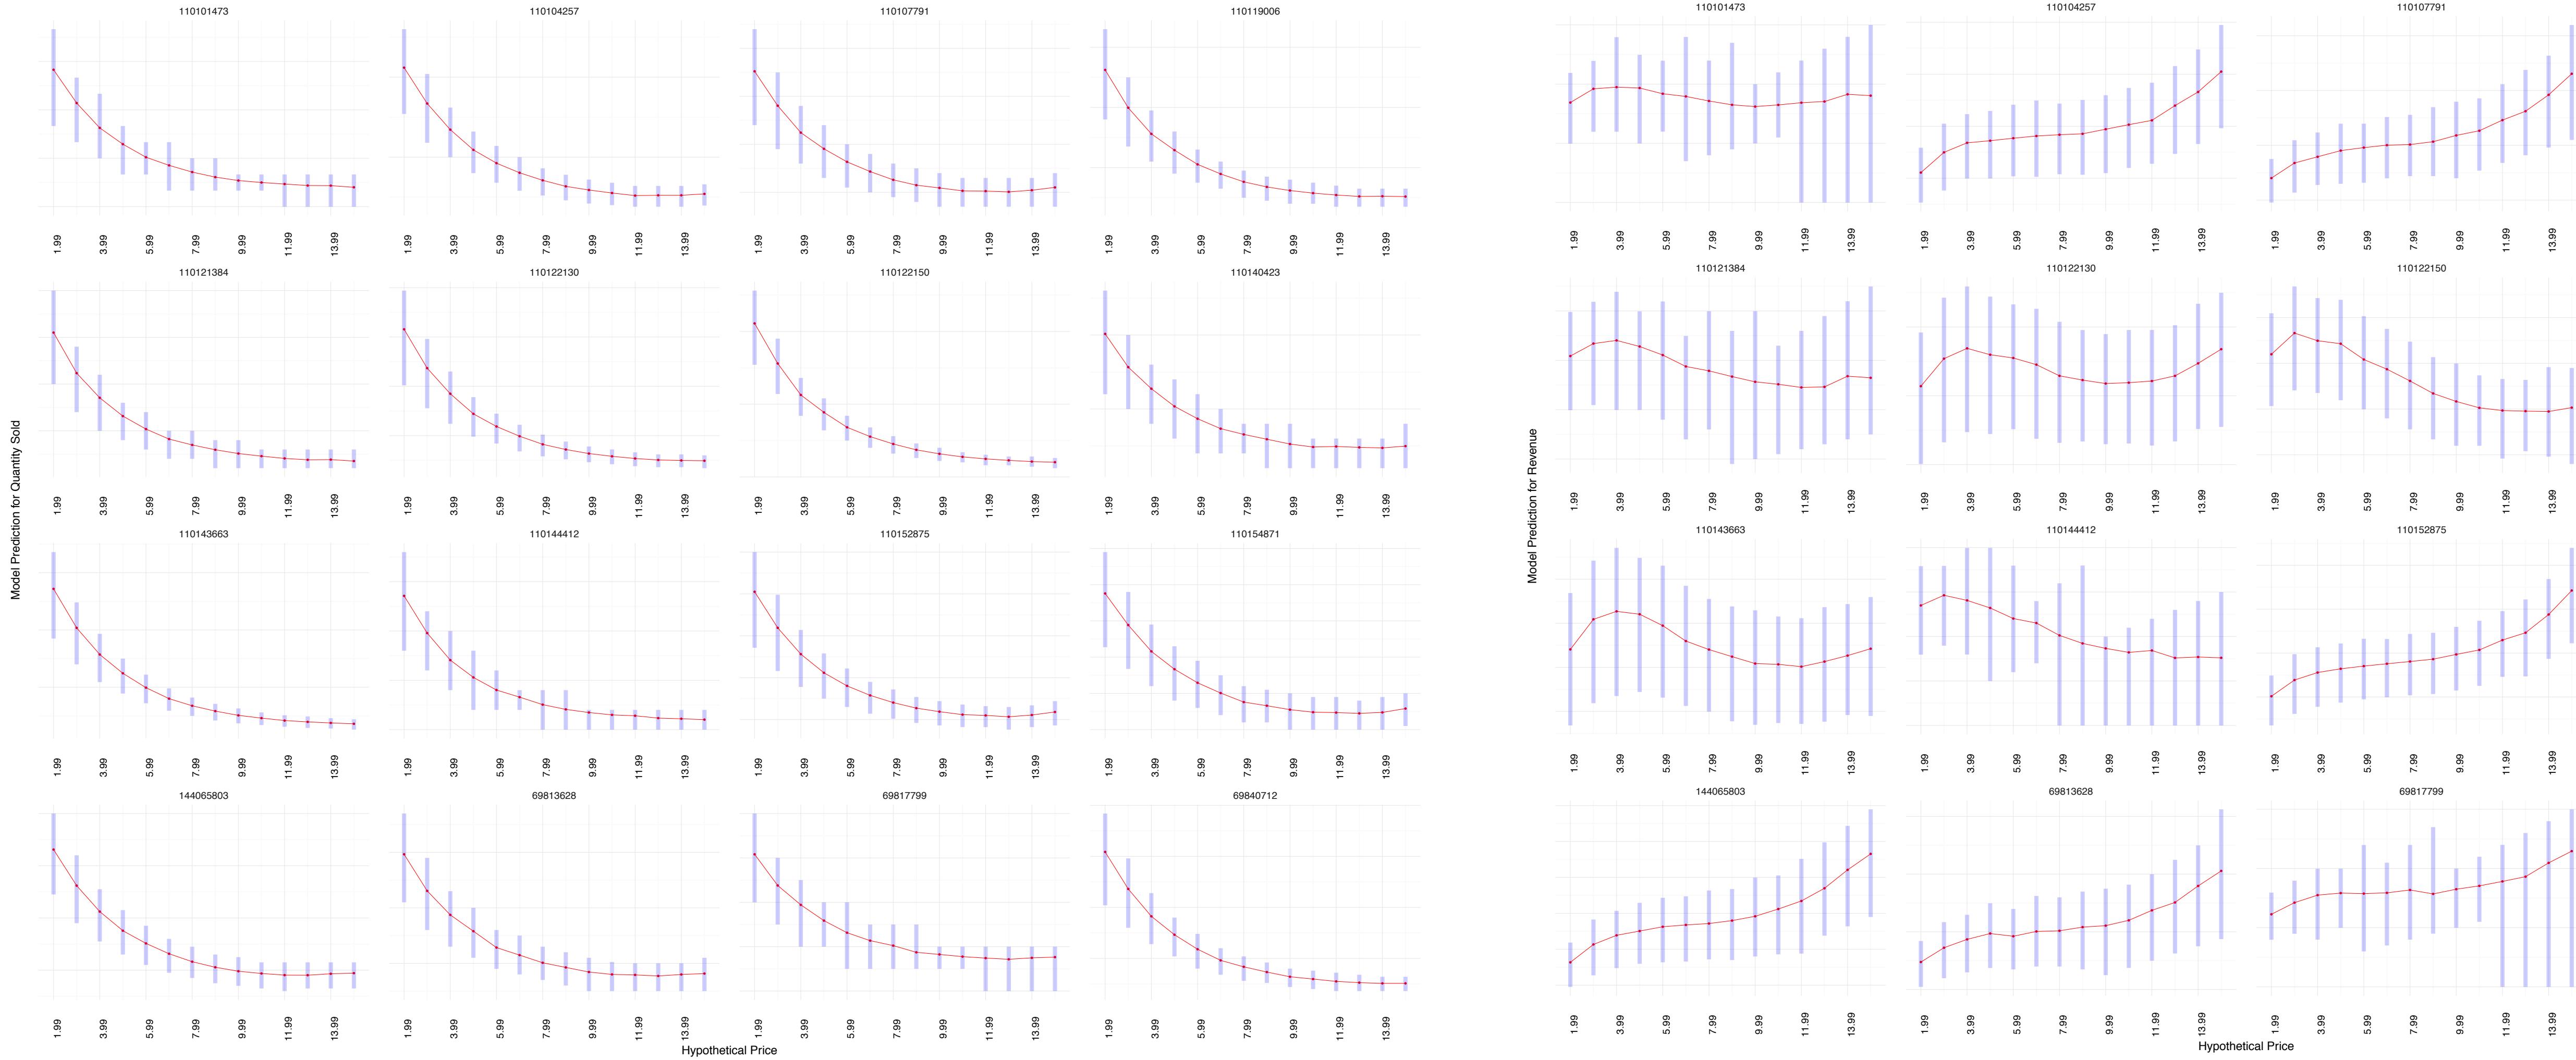
```
# A tibble: 1,946 x 7
```

|    | prod_key_factor | prod_key   | price | ysd_scaled | price_scaled | price_sqr_scaled | month |
|----|-----------------|------------|-------|------------|--------------|------------------|-------|
|    | <fctr>          | <chr>      | <dbl> | <dbl>      | <dbl>        | <dbl>            | <dbl> |
| 1  | aaaaaaaaaa      | aaaaaaaaaa | 1.99  | 1.595587   | -3.58149701  | -2.5113317       | 8     |
| 2  | aaaaaaaaaa      | aaaaaaaaaa | 2.99  | 1.595587   | -3.19446355  | -2.4144849       | 8     |
| 3  | aaaaaaaaaa      | aaaaaaaaaa | 3.99  | 1.595587   | -2.80743009  | -2.2787437       | 8     |
| 4  | aaaaaaaaaa      | aaaaaaaaaa | 4.99  | 1.595587   | -2.42039662  | -2.1041082       | 8     |
| 5  | aaaaaaaaaa      | aaaaaaaaaa | 5.99  | 1.595587   | -2.03336316  | -1.8905783       | 8     |
| 6  | aaaaaaaaaa      | aaaaaaaaaa | 6.99  | 1.595587   | -1.64632970  | -1.6381541       | 8     |
| 7  | aaaaaaaaaa      | aaaaaaaaaa | 7.99  | 1.595587   | -1.25929624  | -1.3468357       | 8     |
| 8  | aaaaaaaaaa      | aaaaaaaaaa | 8.99  | 1.595587   | -0.87226277  | -1.0166228       | 8     |
| 9  | aaaaaaaaaa      | aaaaaaaaaa | 9.99  | 1.595587   | -0.48522931  | -0.6475157       | 8     |
| 10 | aaaaaaaaaa      | aaaaaaaaaa | 10.99 | 1.595587   | -0.09819585  | -0.2395142       | 8     |

```
# ... with 1,936 more rows
```

```
pred_q <- posterior_predict(fit, newdata = new_data)
```

# Computing Demand Curves and Revenue Predictions

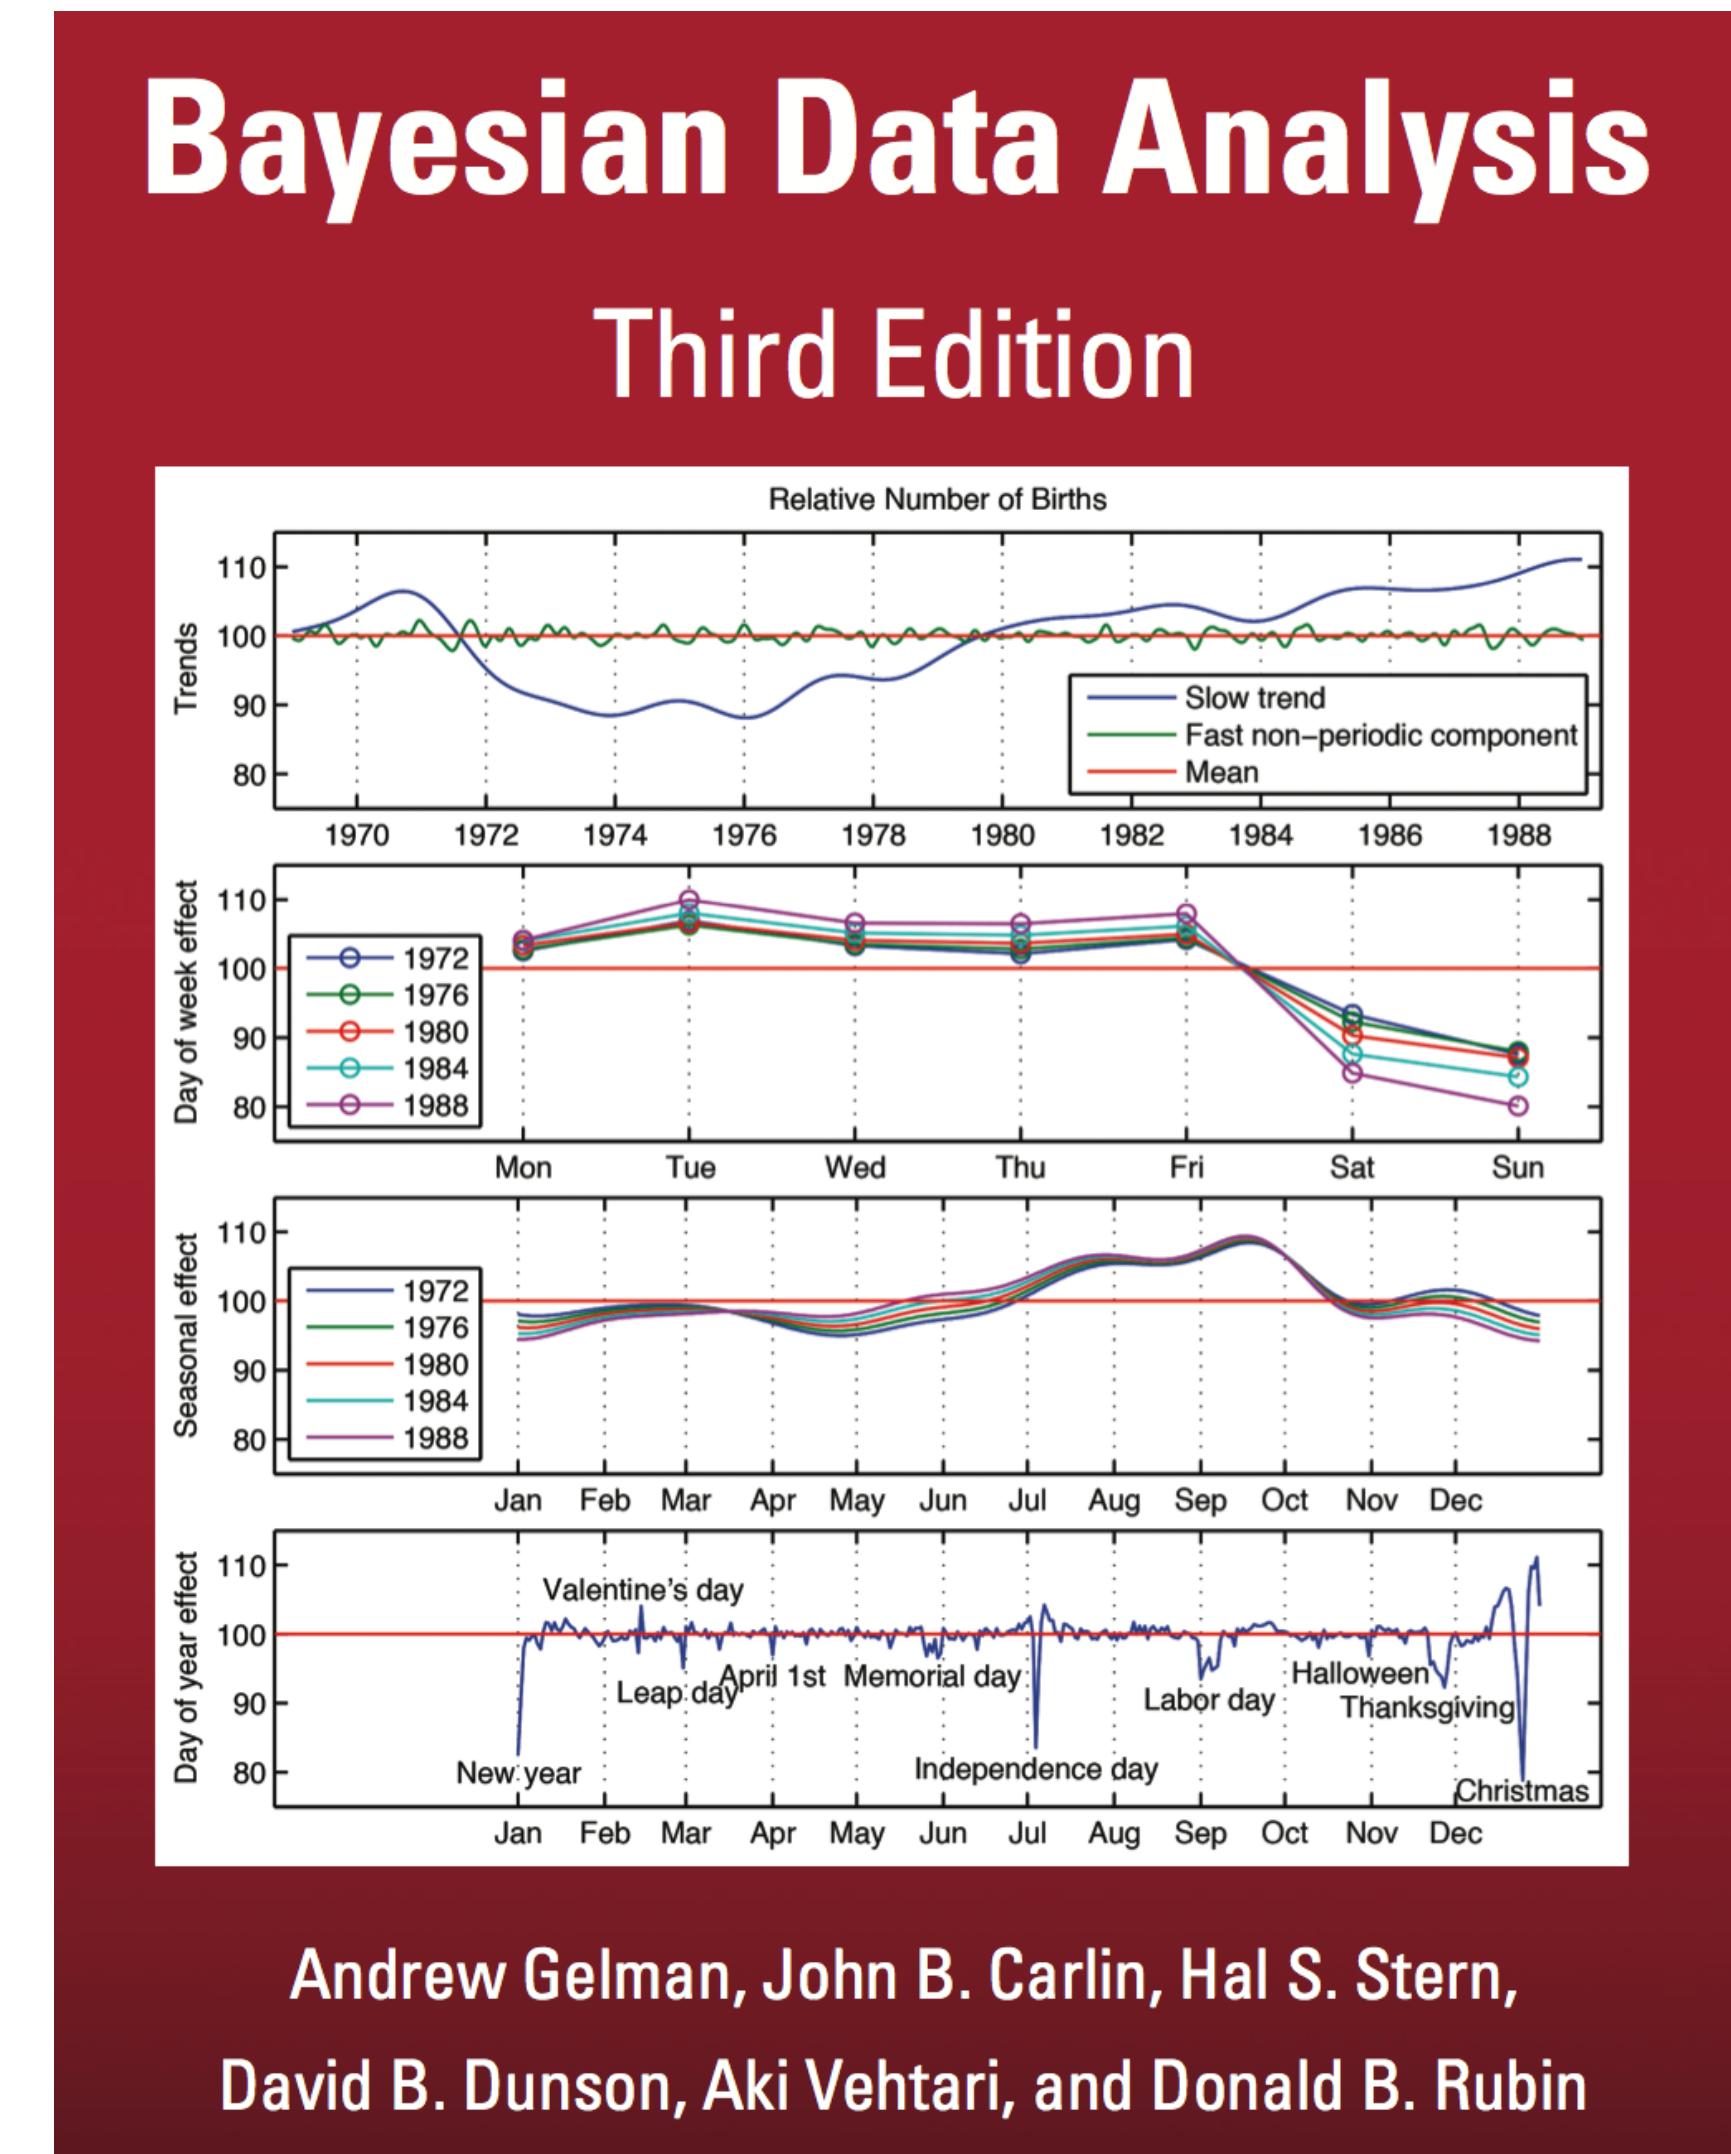
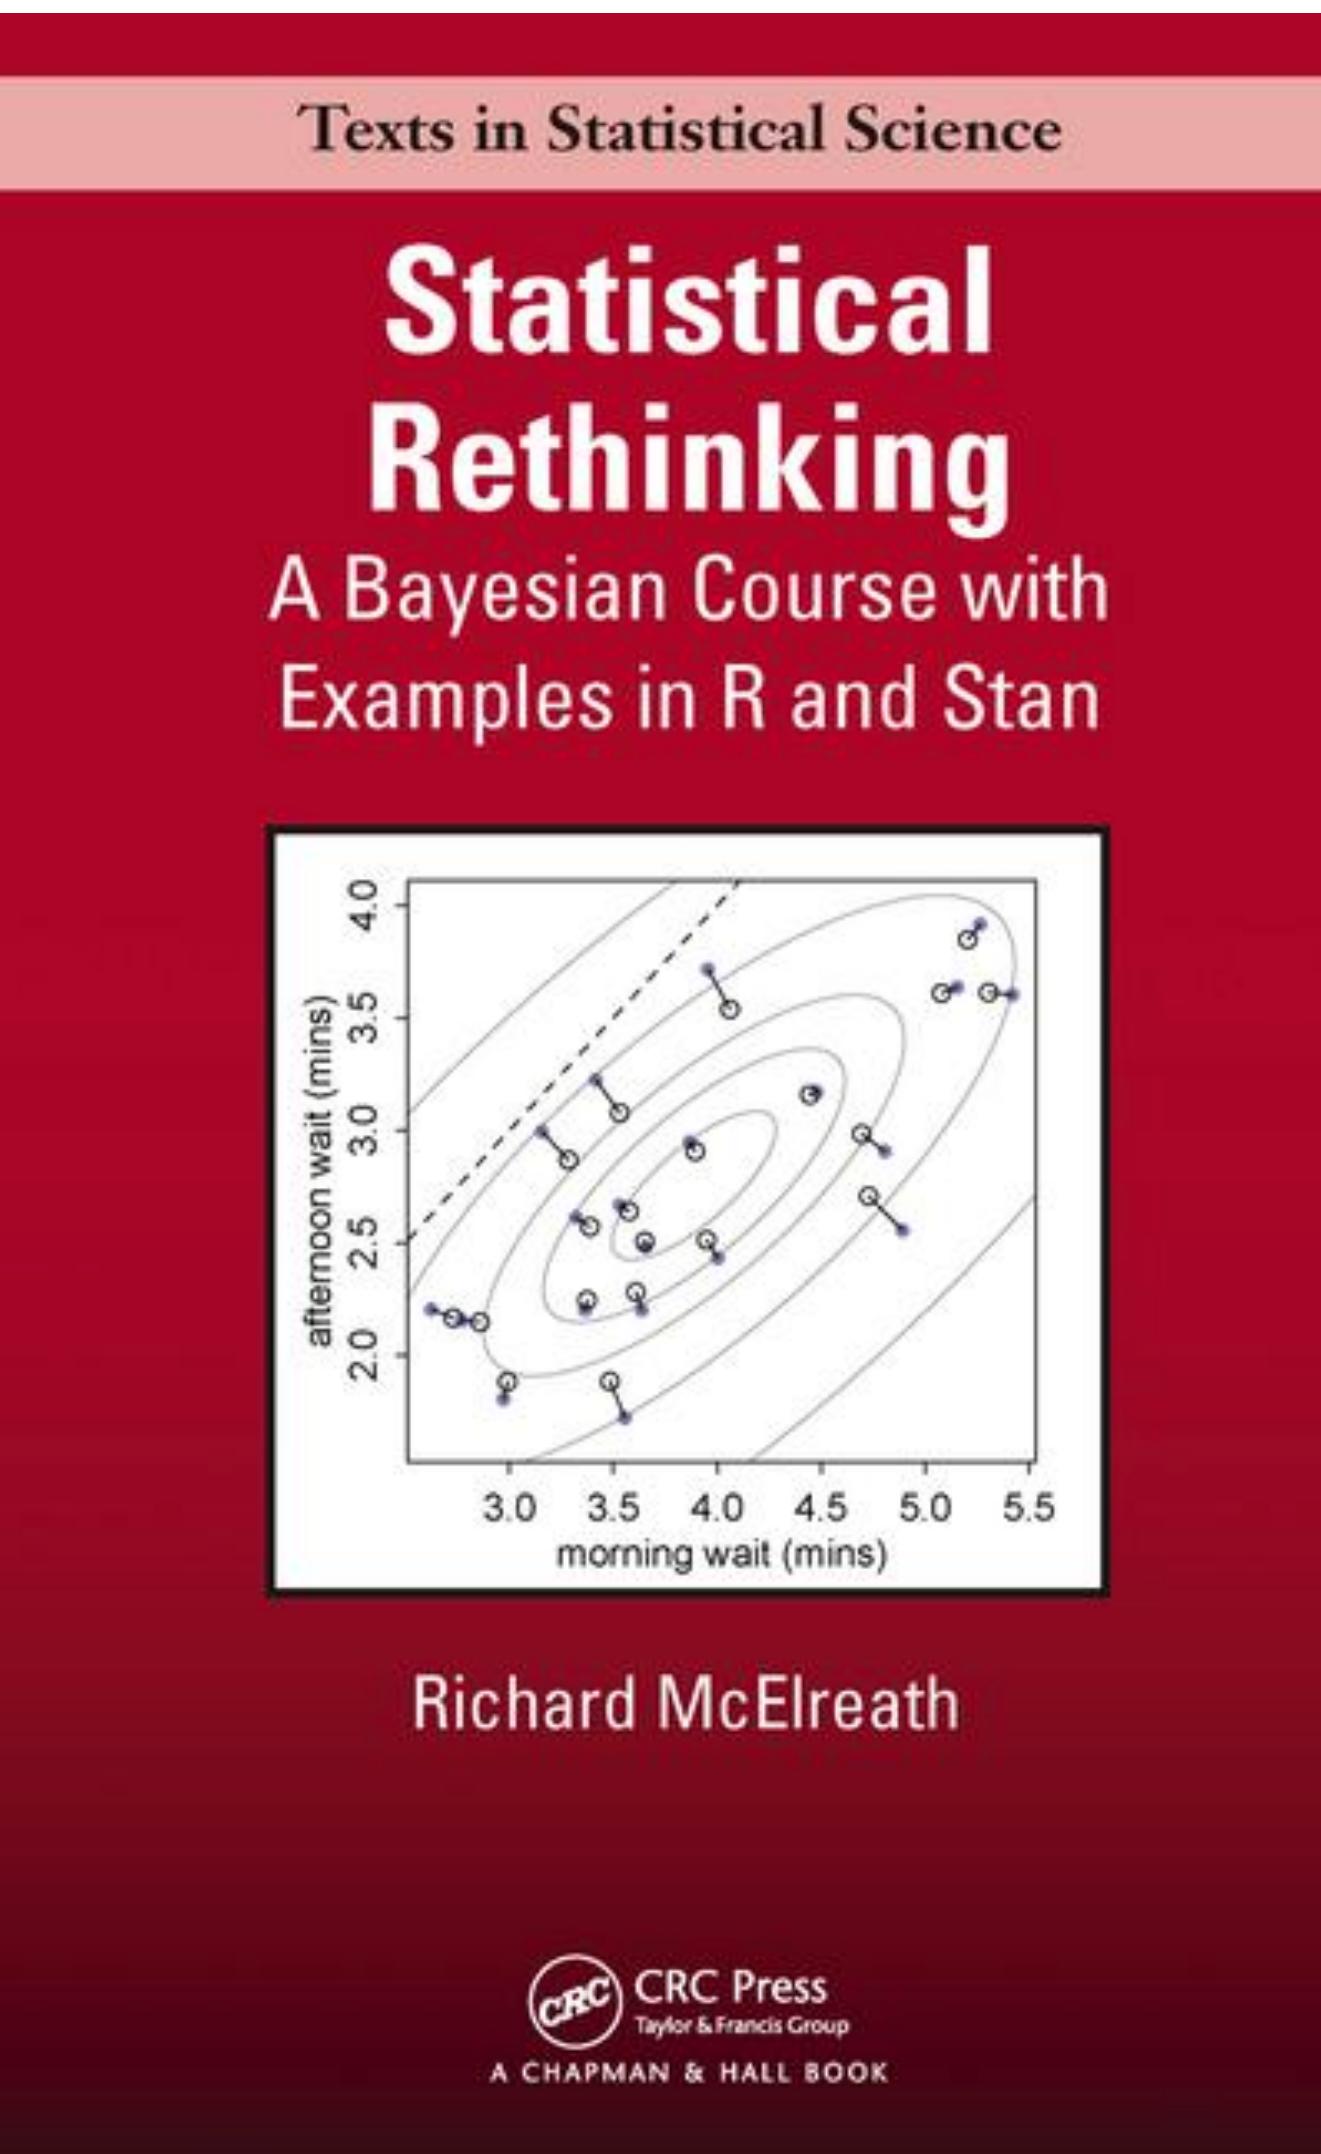
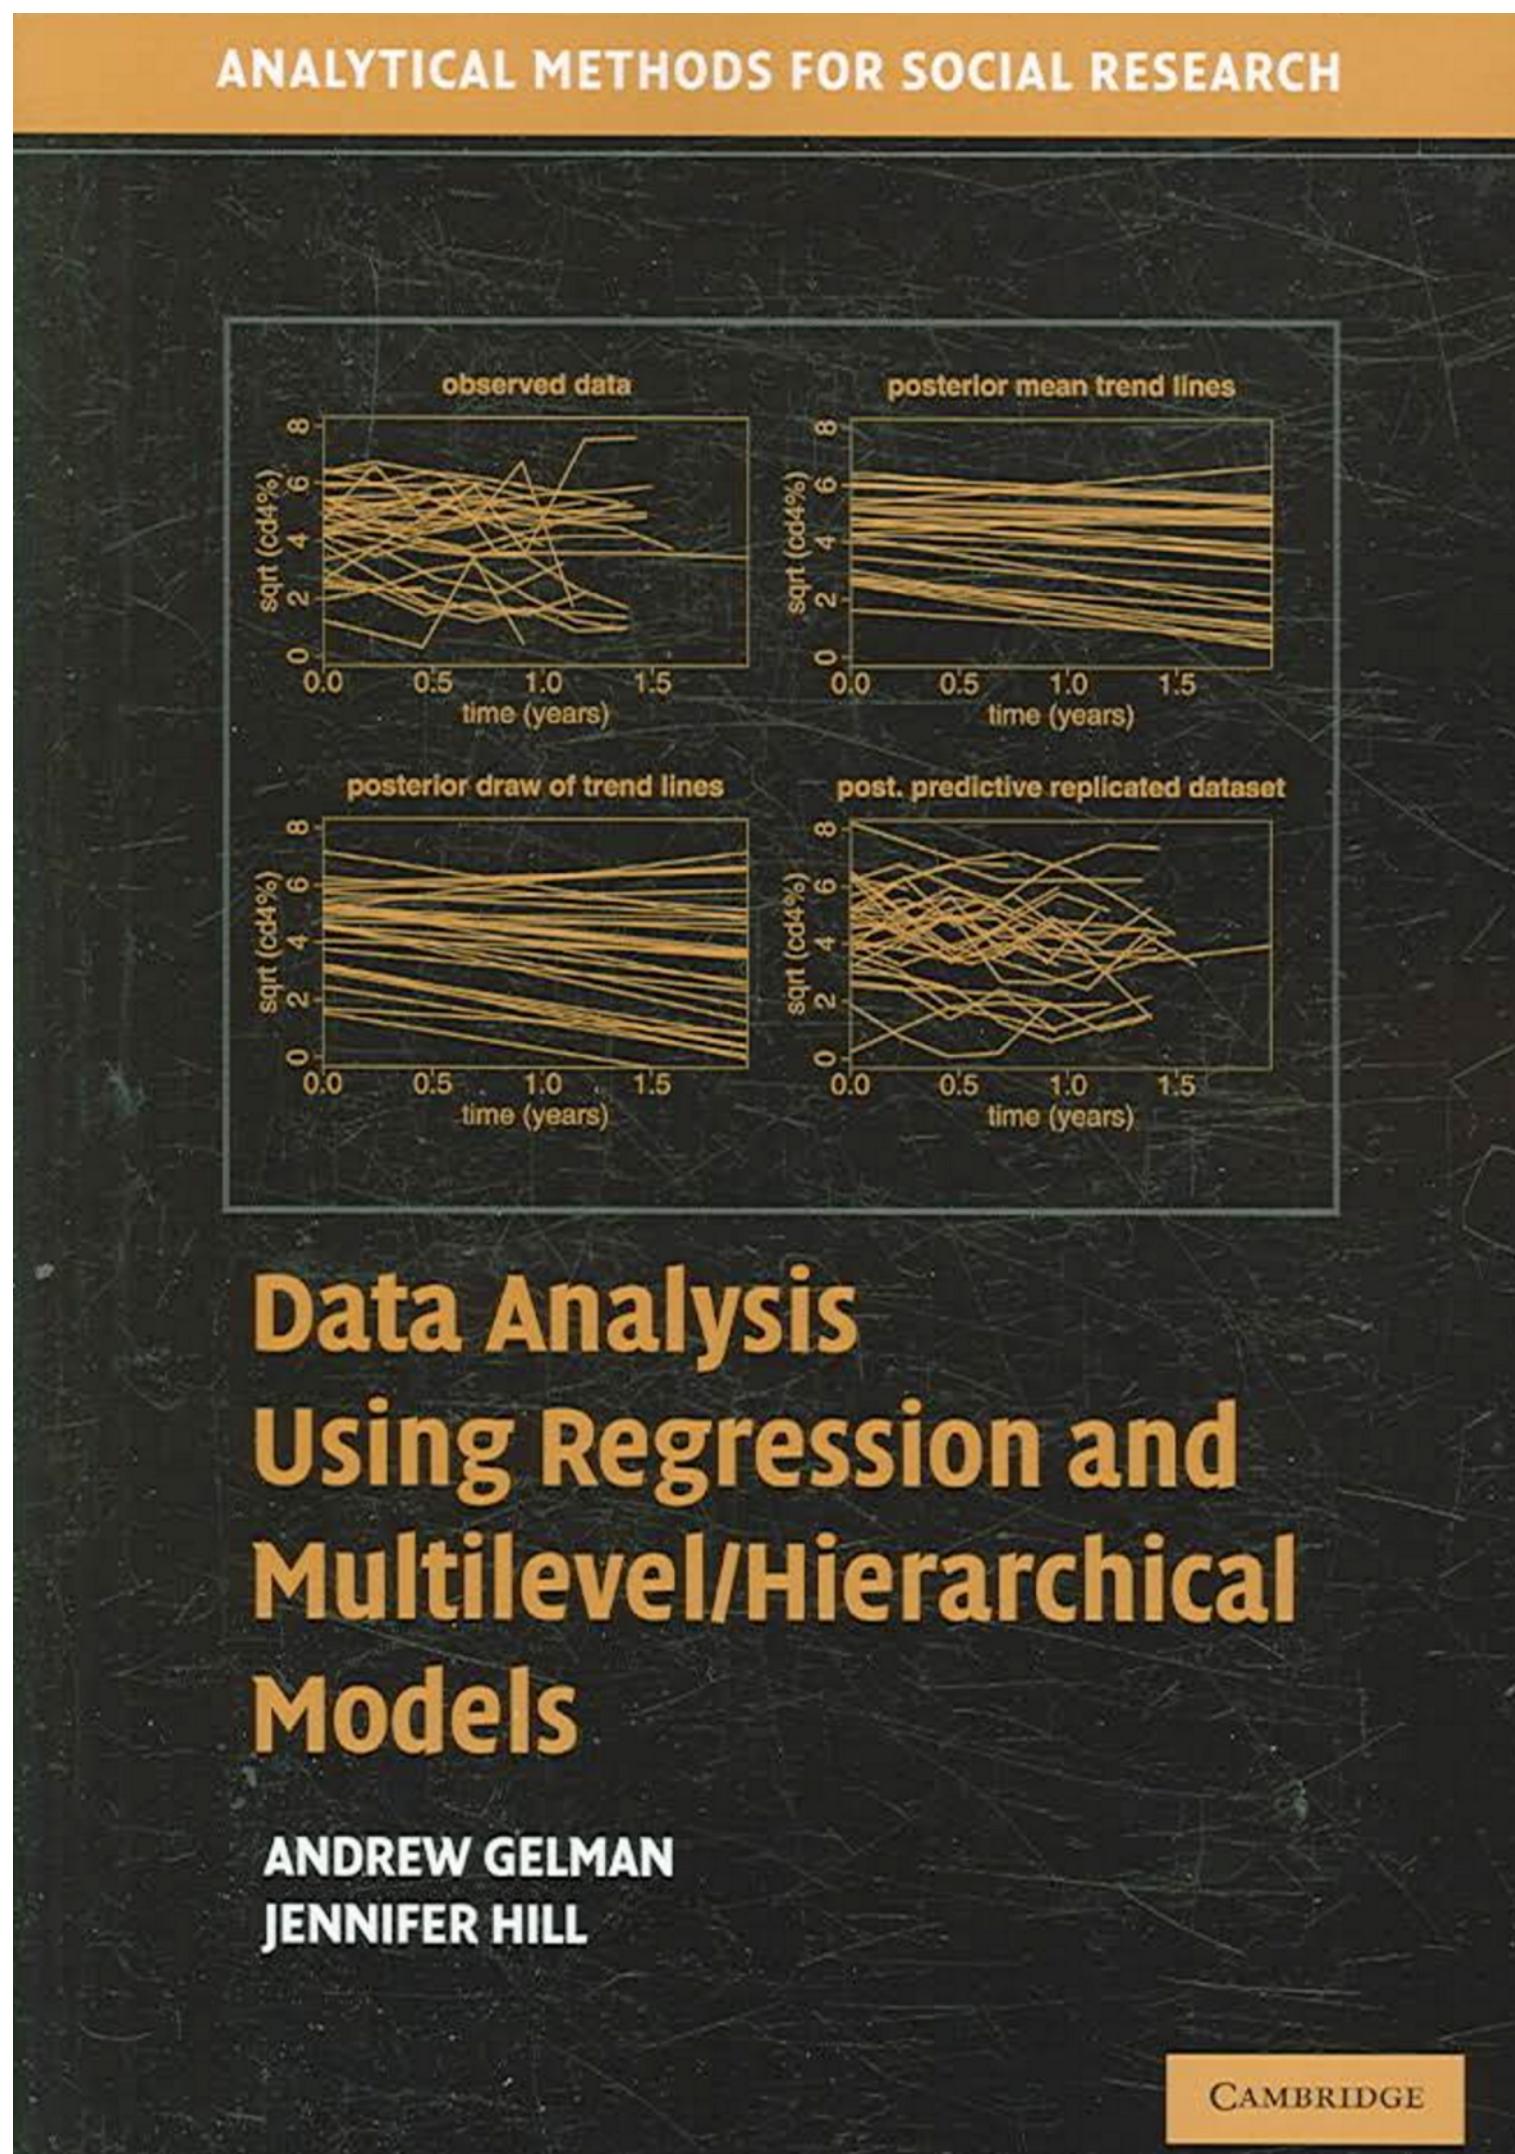


# References

---



# Books



# Some Papers and Videos

---

- ▶ Stan: A probabilistic programming language for Bayesian inference and optimization (Andrew Gelman, et. al.) [http://www.stat.columbia.edu/~gelman/research/published/stan\\_jebs\\_2.pdf](http://www.stat.columbia.edu/~gelman/research/published/stan_jebs_2.pdf)
- ▶ Stan: A Probabilistic Programming Language (Bob Carpenter, et. al.) <http://www.stat.columbia.edu/~gelman/research/published/stan-paper-revision-feb2015.pdf>
- ▶ Hamiltonian Monte Carlo (Michael Betancourt) <https://www.youtube.com/watch?v=pHsuIaPbNbY>
- ▶ Stan Hands-on with Bob Carpenter <https://www.youtube.com/watch?v=6NXRCtWQNMg>
- ▶ A lot more available on [mc-stan.org](http://mc-stan.org)

# Merci Beaucoup!

- ▶ eric@stan.fit
- ▶ @ericnovik
- ▶ mc-stan.org / stan.fit
- ▶ [www.linkedin.com/in/enovik](https://www.linkedin.com/in/enovik)

